

R PROFILING AND OPTIMISATION

GÜNTHER SAWITZKI

PENDING CHANGES

Warning: this is under construction.

This vignette contains experimental material which may sink down to the package implementation, or vanish.

Known issues:

- Control information may be included as special stack in raw format.
- A list of profiles may become default. Only one profiling interval value per profile.
- Nodes may be implemented as *factor*. Work-around for the R factor handling needs to be added, i.e. *factor* as a data structure.
- changing timing interval is too expensive, as rle is not transparent to data frames. Implement profiles as a list, with a time interval attribute per list element.

CONTENTS

Pending changes	1
Profiling facilities in R	2
L ^A T _E X layout tools and R settings	3
1. Profiling	5
1.1. Simple regression example	6
1.1.1. R basic	6
1.1.2. Package sprof	9
1.1.3. Node classes	13
2. A better grip on profile information	20
2.1. The internal details	20
2.2. The free lunch	25
2.3. Cheap thrills	26
2.3.1. Trimming	30
2.3.2. Surgery	34

Date: May 2013. *Revised:* July 2013

Typeset, with minor revisions: July 30, 2013 from SVN *Revision* : 187 2013-07-30.

Key words and phrases. R programming, profiling, optimisation, R program language.

An R vignette for package sprof.

URL: <http://sintro.r-forge.r-project.org/>

Private Version

2.4. Run length	36
3. xxx – lost & found	45
3.1. Graph package	53
4. Standard output	55
4.1. Print	55
4.2. Summary	55
4.3. Plot	56
5. Graph	56
5.1. Example: regression	56
5.1.1. graph package	57
5.1.2. igraph package	58
5.1.3. network package	59
5.1.4. Rgraphviz package	61
5.2. Trimmed example: regression	66
5.2.1. graph package	67
5.2.2. igraph package	70
5.2.3. network package	71
5.2.4. Rgraphviz package	76
Index	81

PROFILING FACILITIES IN R

R provides the basic instruments for profiling, both for time based samplers as for event based instrumentation. Information on R profiling is in Section 3.2 “Profiling R code for speed” and section “3.3 Profiling R code for memory use”. of Writing R Extensions <http://cran.r-project.org/doc/manuals/R-exts.html>. Specific information on memory profiling is in <http://developer.r-project.org/memory-profiling.html>.

However this source of information seems to be rarely used.

Maybe the supporting tools are not adequate. The summaries provided by R reduce the information beyond necessity. Additional packages are available, but these are not sufficiently action oriented.

With package *sprof* we want to give a data representation that keeps the full profile information. Tools to answer common questions are provided. The data structure should make it easy to extend the tools as required.

The package is currently distributed at r-forge as part of the *sintro* material.

To install this package directly within R, type

```
install.packages("sprof", repos="http://r-forge.r-project.org")
```

To install the recent package from source directly within R, type

```
install.packages("sprof", repos="http://r-forge.r-project.org", type="source")
```

L^AT_EX LAYOUT TOOLS AND R SETTINGS

You may want to skip this section, unless you want to modify the vignette for your own purposes, or look at the internals.

This is the main library we are going to use.

Input

```
library(sprof)
```

We want immediate warnings, if necessary. Set to level 2 to handle warnings as error.

Input

```
message("switching options(warn=1) -- immediate warning on")
options(warn=1)
```

We want a second chance on errors. So install an error handler.

Input

```
options(error = recover)
```

Print parameters used here:

Input

```
options(width = 72)
options(digits = 6)
```

For *str* output, we generally use these settings:

Input

```
strx <- function(x, ...){
  str(x, max.level=2, vec.len=3, nchar.max=40,
    list.len=12,
    width=70, strict.width="wrap",...)
}
```

For larger tables and data frames, we use a kludge to avoid long outputs.

ToDo: add `keep3` to
keep header, some
middle, tail

Input

```
xcutdata.frame <- function(df, cut, margin){
  #! keep3, to add: margin top - random center - margin bottom
  if (!is.data.frame(df)) return(df)
  nrow <- nrow(df)
  # cut a range if it is not empty.
  # Quiet noop else.
  # Does not cut single lines.
  cutrng <- function(cutfrom,cutto){
    if (cutfrom<cutto){
      df[cutfrom,] <- NA
      if (!is.null(rownames(df))) rownames(df)[cutfrom] <- "< cut >"
      if (!is.null(df$name)) df$name[cutfrom] <- ""

      cutfrom <- cutfrom+1
    }
  }
}
```

```

        df[-(cutfrom:cutto),]
      }#if
    }
    if (!missing(cut)) {df <- cutrng(cut[1],cut[2]); return(df)}
    if (!missing(margin)) {
      if (length(margin)==1) margin <- c(margin,margin)
      cut <- c(margin[1]+1,nrow-margin[2])
      df <-cutrng(cut[1],cut[2]);
      return(df)}
    #   if (!missing(keep3)) { cut <- c(keep3[1]+1, keep3[1]+1,
    #                                   nrow-keep3[3]-1,nrow-keep3[3]-1)
    #   if (cut[3]-cut[4] > keep3[2]+2){delta<-(cut[3]-cut[2]) div 2
    #   cut[3]<-0
    #   browser()
    #   } else df <- cutrng(cut[1],cut[4])
    #   cutrng(cut[1],cut[4]) return(df)}
  }

```

ToDo: remove text
vdots from string/-
name columns. Use
empty string.

We use the R function `xtable()` for output and \LaTeX `longtable`. A convenient wrapper to use this in our `Sweave` source is:

```

library(xtable)
prxt <- function (x, digits=2, cut=TRUE, caption=NULL,
  label=NULL, ...) {
  if (cut) {margin <- 10
  if (nrow(x)> 2*margin+3) x <-xcutdata.frame(x, margin=margin)}
  print(
    xtable(x, digits=digits, caption=caption,
            label=label, ...),
    floating=FALSE,
    tabular.environment="longtable",
    caption.placement="top",
    zero.print = ".",
    NA.string="\vdots")
  }

```

This is to be used with `<<print=FALSE, results =tex, label=tab:prxx>>=`

The graph visualisation family is not friendly. We try to get control by using a wrapper which is at least used to the members of the `graphviz` clan. This will be used in later sections.

```

plotviz <- function(x,...)
{
  plot(x,
    attrs=list(node=list(cex=4, fontsize=40, shape="ellipse")),
    cex.main=2, ...)
}

```

1. PROFILING

The basic information provided by all profilers is a protocol of sampled stacks. For each recorded event, the protocol has one record, such as a line with a text string showing the sampled stack.

We use profiles to provide hints on the dynamic behaviour of programs. Most often, this is used to improve or even optimise programs. Sometimes, it is even used to understand some algorithm.

Profiles represent the program flow, which is considered to be laid out by the control structure of a program. The control structure is represented by the control graph, and this leads to the common approach to (re)construct the control graph, map the profile to this graph, and used graph based methods for further analysis. The prime example for this strategy is the GNU profiler *gprof* (see <http://sourceware.org/binutils/docs/gprof/>) which is used as master plan for many common profilers.

It is only half of the truth that the control graph can serve as a base for the profiled stacks. In R, we have some peculiarities.

lazy evaluation: Arguments to functions can be passed as promises. These are only evaluated when needed, which may be at a later time, and may then lead to insertions in the stack. So we may have information resulting from the data flow, interspersed with the control flow.

memory management: Allocation of memory, and garbage collection, may interfere and leave their traces in the stack. While allocation is closely related to the visible control flow, garbage collection is a collective effect largely out of control of the code to execute.

primitives: Internal functions may escape the usual stack conventions and execute without leaving any identifiable trace on the stack.

control structures: In R, many control structures are implemented as function. Most notably, the *apply()* family appears as function calls and can lead to cliques in the graph representation that do not correspond to relevant structures. Since these functions are well known, they can have a special treatment.

So while the stack follows an overall well known dynamics, in R there are exceptions from regularity.

The general approach, by *summaryRprof()* and others, is to reduce the profile to node information, or to consider single transitions.

We take a different approach. We take the stacks, as recorded in the profiles as our basic information unit. From this, we ask: what are the actions we need to answer our questions? Representation in graphs may come later, if they can help.

If the stacks would come from the control flow only, we could make use of the sequential nature of stacks. But since we have to live with the R specific interferences, we stay with the raw stacks.

In this presentation, we will use a small list of examples. Since *Rprof* is not implemented on all systems, and since the profiles tend to get very large, we use

ToDo: rearrange
stacks? detect
order?

some prepared examples that are frozen in this vignette and not included in the distribution, but all the code to generate the examples is provided.

1.1. Simple regression example.

Input

```
n <- 10000
x <- runif(n)
err <- rnorm(n)
y <- 2+ 3 * x + err
reg0data <- data.frame(x=x, y=y, err=err)
rm(x,y,err)
```

We will use this example to illustrate the basics. Of course the immediate questions are the variance between varying samples, and the influence of the sample size n . We keep everything fixed, so the only issue for now is the computational performance under strict iid conditions.

Still we have parameters to choose. We can determine the profiling granularity by setting the timing interval, and we can use repeated measurements to increase precision below the timing interval.

The timing interval should depend on the clock speed. Using for example 1ms amounts to some 1000 steps on a current CPU, per kernel.

If we use repeated samples, the usual rules of statistics applies. So taking 100 runs and taking the mean reduces the standard deviation by a factor 1/10.

By the usual R conventions, seconds are used as time base for parameters. However report will use ms as a time base.

Here is an example how to take a profile, using basic R. See section 1.1.2 on page 9 how to use *sampleRprof* in package *sprof* for an easier solution.

Input

```
profininterval <- 0.001
simruns <- 100
Rprof(filename="RprofsRegressionExpl.out", interval = profinterval)
  for (i in 1:simruns) xxx<- summary(lm(y~x, data=reg0data))
Rprof(NULL)
```

We now have the profile data in a file *RprofsRegressionExpl.out*. For this vignette, we use a frozen version *RprofsRegressionExpl01.out*.

1.1.1. *R basic*. The basic R functions invite us to get a summary.

Input

```
sumRprofRegressionExpl <- summaryRprof("RprofsRegressionExpl01.out")
#str(profile_nodes_rle, max.level=2, vec.len=3, nchar.max=40, list.len=6)
strx(sumRprofRegressionExpl)
```

Output

ToDo: Can we calibrate times to CPU rate? Introduce cpu clock cycle as a time base

```

List of 4
$ by.self : 'data.frame': 41 obs. of 4 variables:
..$ self.time : num [1:41] 0.087 0.057 0.051 0.043 0.042 0.04 0.032
  0.026 ...
..$ self.pct : num [1:41] 16.67 10.92 9.77 8.24 ...
..$ total.time: num [1:41] 0.113 0.099 0.069 0.043 0.474 0.045 0.033
  0.114 ...
..$ total.pct : num [1:41] 21.65 18.97 13.22 8.24 ...
$ by.total : 'data.frame': 62 obs. of 4 variables:
..$ total.time: num [1:62] 0.522 0.522 0.521 0.521 0.521 0.521 0.521
  0.521 ...
..$ total.pct : num [1:62] 100 100 99.8 99.8 ...
..$ self.time : num [1:62] 0.006 0 0.001 0 0 0 0 0 ...
..$ self.pct : num [1:62] 1.15 0 0.19 0 0 0 0 0 ...
$ sample.interval: num 0.001
$ sampling.time : num 0.522

```

The summary reduces the information contained in the profile to marginal statistics per node. This is provided in two data frames giving the same information, only in different order.

The file contains several spurious recordings: nodes that have been recorded only few times. It is worth noting these, but then they better be discarded. We use a time limit of 4ms, which given our sampling interval of 1ms means we require more than four observations.

```

prxt(sumRprofRegressionExpl$by.self,
      caption="summaryRprof result: by.self as final stack entry, all records",
      label="tab:prSRREbs")

```

Table 1: summaryRprof result: by.self as final stack entry, all records

	self.time	self.pct	total.time	total.pct
"lm.fit"	0.09	16.67	0.11	21.65
"[.data.frame"	0.06	10.92	0.10	18.97
"model.matrix.default"	0.05	9.77	0.07	13.22
"as.character"	0.04	8.24	0.04	8.24
"lm"	0.04	8.05	0.47	90.80
"summary.lm"	0.04	7.66	0.04	8.62
"structure"	0.03	6.13	0.03	6.32
"na.omit.data.frame"	0.03	4.98	0.11	21.84
"anyDuplicated.default"	0.02	4.21	0.02	4.21
"as.list.data.frame"	0.02	4.21	0.02	4.21
< cut >	:	:	:	:
"FUN"	0.00	0.19	0.01	1.34
"%in%"	0.00	0.19	0.00	0.77
"deparse"	0.00	0.19	0.00	0.38
"\$"	0.00	0.19	0.00	0.19
"as.list.default"	0.00	0.19	0.00	0.19

"as.name"	0.00	0.19	0.00	0.19
"coef"	0.00	0.19	0.00	0.19
"file"	0.00	0.19	0.00	0.19
"NCOL"	0.00	0.19	0.00	0.19
"terms.formula"	0.00	0.19	0.00	0.19

```

Input
prxt(sumRprofRegressionExpl$by.total[sumRprofRegressionExpl$by.total$total.time>0.004,],
caption="summaryRprof result: by.total, total time > 4ms",
label="tab:prSRREbt")

```

Table 2: summaryRprof result: by.total, total time > 4ms

	total.time	total.pct	self.time	self.pct
"<Anonymous>"	0.52	100.00	0.01	1.15
"Sweave"	0.52	100.00	0.00	0.00
"eval"	0.52	99.81	0.00	0.19
"doTryCatch"	0.52	99.81	0.00	0.00
"evalFunc"	0.52	99.81	0.00	0.00
"try"	0.52	99.81	0.00	0.00
"tryCatch"	0.52	99.81	0.00	0.00
"tryCatchList"	0.52	99.81	0.00	0.00
"tryCatchOne"	0.52	99.81	0.00	0.00
"withVisible"	0.52	99.81	0.00	0.00
< cut >	:	:	:	:
"as.list"	0.02	4.41	0.00	0.00
"anyDuplicated.default"	0.02	4.21	0.02	4.21
"as.list.data.frame"	0.02	4.21	0.02	4.21
"sapply"	0.01	2.68	0.00	0.19
"match"	0.01	2.11	0.00	0.19
"[[.data.frame"	0.01	1.53	0.00	0.19
"["	0.01	1.53	0.00	0.00
"rep.int"	0.01	1.34	0.01	1.34
"FUN"	0.01	1.34	0.00	0.19
"list"	0.01	0.96	0.01	0.96

1.1.2. *Package sprof*. In contrast to the common R packages, in the *sprof* implementation we take a two step approach. First we read in the profile file to an internal representation. Analysis is done in later steps.

Input

```
sprof01<- readRprof("RprofsRegressionExpl01.out")
```

The data contain identification information for reference. This will be used in the functions of *sprof* and shown in the displays. Here is the summary of this section:

Input

```
str(sprof01$info)
```

Output

```
'data.frame':      1 obs. of  8 variables:
 $ id      : Factor w/ 1 level  "\"RprofsRegressionExpl01.out\" 2013-06-13 23:46:04": 1
 $ date    : POSIXct, format: "2013-07-30 15:17:11"
 $ nrnodes : int 62
 $ nrstacks: int 50
 $ nrrecords: int 522
 $ firstline: Factor w/ 1 level "sample.interval=1000": 1
 $ ctllines : Factor w/ 1 level "sample.interval=1000": 1
 $ ctllinenr: num 1
```

For this vignette, we change the *id* information. So in this context:

Input

```
sprof01$info$id <- "sprof01"
```

We keep this example and use the copy *sprof01* of it extensively for illustration.

Input

```
save(sprof01, file="sprof01lm.RData")
```

To run the vignette with a different profile, replace *sprof01* by your example. You still have the file for reference.

Package *sprof* provides a function *sampleRprof()* to take a sample and create a profile on the fly, as in

Input

```
sprof01temp <- sampleRprof(runif(10000), runs=100)
```

The basic data structure consists of four data frames. The *info* section collects global information from the input file, such as an identification strings and various global matrix. The *nodes* section initially gives the same information marginal information as *summaryRprof*. The *stacks* section puts the node information into their calling context as found in the input profile file. The *profiles* section gives the temporal context. It is implemented as a list, but conceptually it is a data frame. Implementing it as a list allows run length encoding of variables, which unfortunately is not allowed by R in data frames.

ToDo: add sampling.interval, sampling.time for backward compatibility

Input

```
strx(sprof01)
```

```

List of 4
$ info : 'data.frame': 1 obs. of 8 variables:
..$ id : chr "sprof01"
..$ date : POSIXct[1:1], format: "2013-07-30 15:17:11"
..$ nrnodes : int 62
..$ nrstacks : int 50
..$ nrrecords: int 522
..$ firstline: Factor w/ 1 level "sample.interval=1000": 1
..$ ctllines : Factor w/ 1 level "sample.interval=1000": 1
..$ ctllinenr: num 1
$ nodes : 'data.frame': 62 obs. of 5 variables:
..$ name : Factor w/ 62 levels "!", "..getNamespace", ...: 1 2 3 4 5 6 7
  8 ...
..$ self.time : num [1:62] 2 0 2 0 0 57 0 1 ...
..$ self.pct : num [1:62] 0.38 0 0.38 0 ...
..$ total.time: num [1:62] 2 1 4 26 99 99 8 8 ...
..$ total.pct : num [1:62] 0.03 0.01 0.05 0.34 1.29 1.29 0.1 0.1 ...
$ stacks : 'data.frame': 50 obs. of 7 variables:
..$ nodes :List of 50
.. .. [list output truncated]
..$ shortname : Factor w/ 50 levels
  "S<A>eFttCtCLtCOdTCwVeesleem.m..n.n...[\"| __truncated__,...: 27 17
  19 1 35 36 37 30 ...
..$ refcount : num [1:50] 1 5 26 55 13 43 51 87 ...
..$ stacklength : int [1:50] 19 20 19 21 14 15 15 14 ...
..$ stackheadnodes: int [1:50] 52 52 52 52 52 52 52 52 ...
..$ stackleafnodes: int [1:50] 27 28 41 6 39 14 38 30 ...
..$ stackssrc : Factor w/ 50 levels "!" [.data.frame [
  na.omit.data.frame na."| __truncated__,...: 27 28 39 5 37 13 36 30
  ...
$ profiles:List of 4
..$ data : int [1:522] 1 2 2 3 4 4 5 5 ...
..$ mem : NULL
..$ malloc : NULL
..$ timesRLE:List of 2
.. ..- attr(*, "class")= chr "rle"
- attr(*, "class")= chr [1:2] "sprof" "list"

```

The nodes do not come in a specific order. Access via a permutation vector is preferred. This allows different views on the same data set. For example, table 4 on page 12 uses a permutation by total time, and a selection (compare to table 2 on page 8). The only difference is that *sprof* works uses a millisecond (ms) base, whereas R in general uses seconds as a base.

```

nodes <- sprofo1$nodes[order(sprofo1$nodes$self.time, decreasing=TRUE),]
prxt(nodes[nodes$self.time>4,],
caption="splot result: by.self, self time > 4ms",
label="tab:prspbtself")

```

Table 3: splot result: by.self, self time > 4ms

	name	self.time	self.pct	total.time	total.pct
30	lm.fit	87.00	16.67	113.00	1.47
6	[.data.frame	57.00	10.92	99.00	1.29
38	model.matrix.default	51.00	9.77	69.00	0.90
14	as.character	43.00	8.24	43.00	0.56
29	lm	42.00	8.05	474.00	6.16
51	summary.lm	40.00	7.66	45.00	0.59
49	structure	32.00	6.13	33.00	0.43
41	na.omit.data.frame	26.00	4.98	114.00	1.48
13	anyDuplicated.default	22.00	4.21	22.00	0.29
16	as.list.data.frame	22.00	4.21	22.00	0.29
40	na.omit	20.00	3.83	134.00	1.74
39	model.response	13.00	2.49	56.00	0.73
36	model.frame.default	12.00	2.30	168.00	2.18
46	rep.int	7.00	1.34	7.00	0.09
10	<Anonymous>	6.00	1.15	522.00	6.79
28	list	5.00	0.96	5.00	0.07

At this level, it is helpful to note the expectations, and only then inspect the timing results. Since we are using a linear model, we are not surprised to see functions related to linear models on the top of the list. We may however be surprised to see functions related to data access and to character conversion very high on the list. The sizeable amount of time spent on NA handling is another aspect that is surprising.

```

Input
nodes <- sprof01$nodes[order(sprof01$nodes$total.time, decreasing=TRUE),]
prxt(nodes[nodes$total.time>4,],
caption="splot result: by.total, total time > 4ms",
label="tab:prspbt")

```

Table 4: splot result: by.total, total time > 4ms

	name	self.time	self.pct	total.time	total.pct
10	<Anonymous>	6.00	1.15	522.00	6.79
52	Sweave	0.00	0.00	522.00	6.79
21	doTryCatch	0.00	0.00	521.00	6.78
22	eval	1.00	0.19	521.00	6.78
23	evalFunc	0.00	0.00	521.00	6.78
55	try	0.00	0.00	521.00	6.78
56	tryCatch	0.00	0.00	521.00	6.78
57	tryCatchList	0.00	0.00	521.00	6.78
58	tryCatchOne	0.00	0.00	521.00	6.78
62	withVisible	0.00	0.00	521.00	6.78
< cut >	\vdots	⋮	⋮	⋮	⋮

61	vapply	3.00	0.57	23.00	0.30
13	anyDuplicated.default	22.00	4.21	22.00	0.29
16	as.list.data.frame	22.00	4.21	22.00	0.29
47	sapply	1.00	0.19	14.00	0.18
31	match	1.00	0.19	11.00	0.14
7	[[0.00	0.00	8.00	0.10
8	[[.data.frame	1.00	0.19	8.00	0.10
25	FUN	1.00	0.19	7.00	0.09
46	rep.int	7.00	1.34	7.00	0.09
28	list	5.00	0.96	5.00	0.07

ToDo: remove text
vdots from string/
name columns

Given the sampling structure of the profiles, two aspects are common. The sampling picks up scaffold functions with a high, nearly constant frequency. And the sampling will pick up rare recordings that are near to detection range. The display functions hide these effects by default. In our example, about half of the nodes are cleared by this garbage collector.

Common rearrangements as by total time and by self time are supplied by the display functions.

`plot_nodes()`, for example, currently gives a choice of four displays for nodes, and supports trimming by default. Our profile starts with 62 nodes. The defaults cut off 34 nodes as uninformative, either because they are too rare, or ubiquitous.

See fig. 1 on the next page.

Information in the time scatterplots may sometimes be more accessible when using a logarithmic scale, so this is added.

If you prefer, you can have the bar charts in horizontal layout, giving more space for labels.

See fig. 2 on page 14.

We can add colour. To illustrate this, we encode the frequency of the nodes as colour. As a palette, we choose a heat map here.

ToDo: apply colour
to selection?

ToDo: spread
colour on displayed
part

Input

```
freqrank01 <- rank(-sprof01$nodes$total.time, ties.method="random")
freqrankcol01 <- heat.colors(length(freqrank01))
```

Here is the node view using these choices:

See fig. 2 on page 14.

Colour is considered a volatile attribute. So you may need to pay some attention to keep colour indices (and colour palettes) aligned to your context. You may want to do experiments with colour, trying to find a good solution for your visual preferences. The recommended way is to use some stable colour index (the slot `icol` is reserved for this) and use this as an index to a choice of colour palettes. So `icol` becomes a part of the data structure, and the colour palette to be used is passed as a parameter.

ToDo: improve
colour: support
colour in a structure

Input



The attribute *icol* is a special case which we used above. If present, it will be interpreted as an index to a colour table. For example, we can collect special well known functions in groups.

ToDo: colour by class – redo. Bundle colour index with colour?

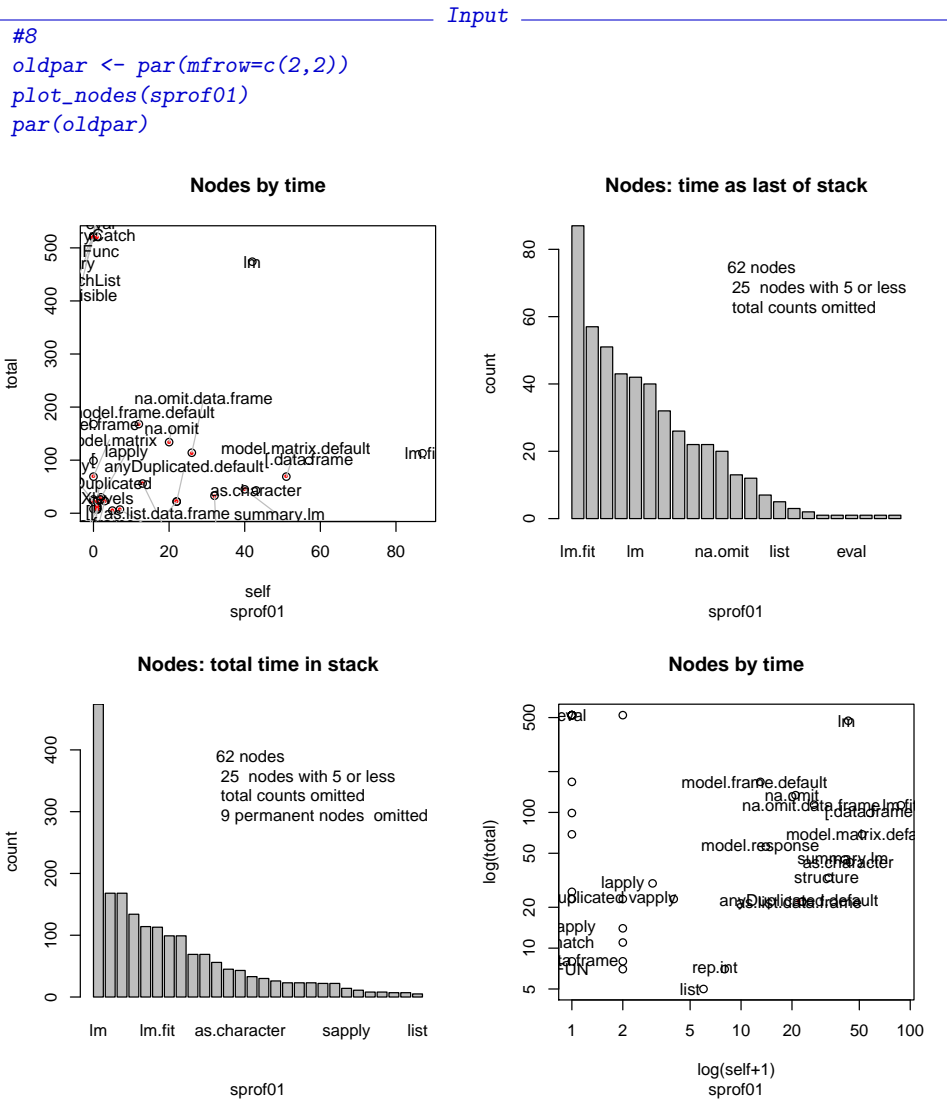


FIGURE 2. Basic information on node level

ToDo: Move class attributes to package code

ToDo: add class by keyword

Grouping may also help you to focus your attention. “HOT” and “cold” may be very helpful tags. These can be used in a flexible way.

```
nodekeyword0 <- function(node)
{
}
```

Input

Input

```
#10
sprof01$nodes$icol <- freqrank01
oldpar <- par(mfrow=c(2,2))
plot_nodes(sprof01, col=freqrankcol01)
par(oldpar)
```

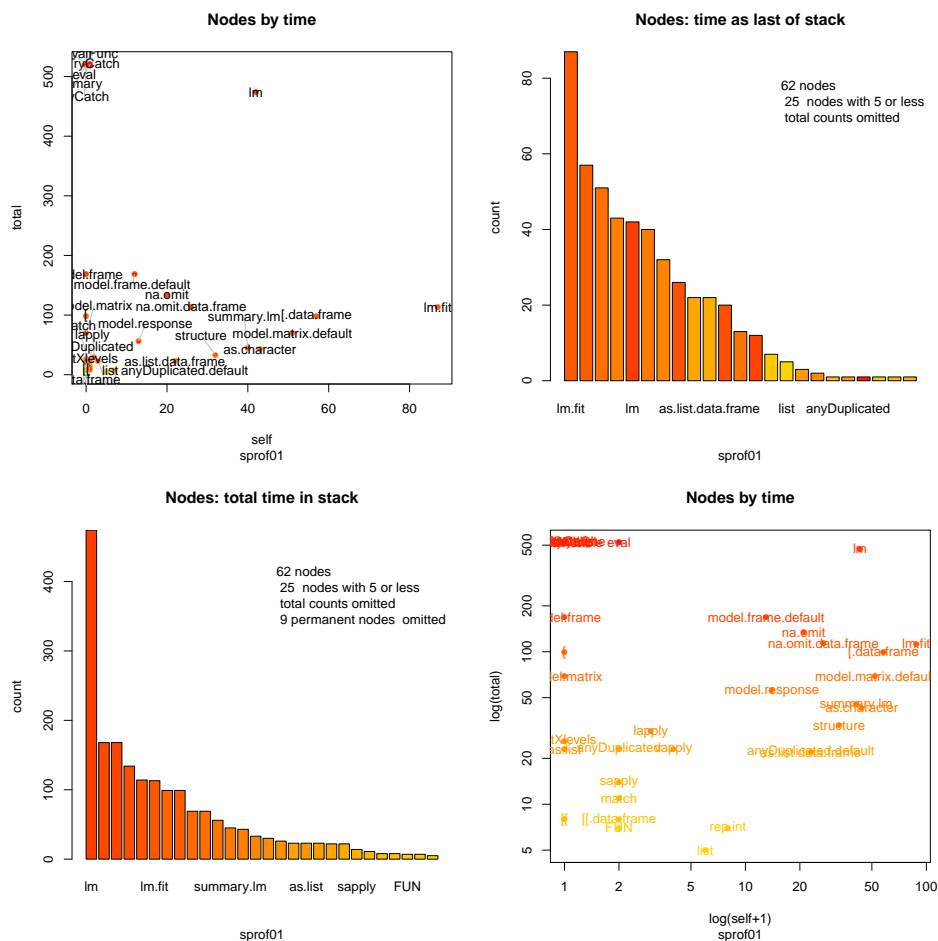


FIGURE 3. Basic information on node level, colour by total time.

```
nodepackages <- nodepackage(sprof01$nodes$name)
names(nodepackages) <- sprof01$nodes$name
table(nodepackages)
```

```
nodepackages
<not found>      base      stats      utils
              6         41         14         1
```

Input

```
sprof01$nodes$icol <-as.factor(nodepackages)
```

See fig. 4 on the facing page.

```
Input
x_apply <- c("apply", "lapply", "vapply", "sapply")
x_as <- c("as.list", "as.data.frame", "as.list.data.frame",
        "as.character", "as.list.default", "as.name")
```

(Extend as you need it) and then use, as for example:

```
Input
nodeclass <- rep("x_nn", sprof01$info$nrnodes)
nodeclass[sprof01$nodes$name %in% x_apply] <- "x_apply"
nodeclass[sprof01$nodes$name %in% x_as] <- "x_as"
```

or use assignments on the fly

```
Input
nodeclass[sprof01$nodes$name %in%
          c("eval", "evalFunc",
            "try", "tryCatch", "tryCatchList", "tryCatchOne",
            "doTryCatch")]
          <- "x_eval"
nodeclass[sprof01$nodes$name %in%
          c("model.frame", "model.matrix.default", "model.frame.default",
            "model.response", "model.matrix", "model.response")]
          <- "x_model"
nodeclass[sprof01$nodes$name %in%
          c("lm", "lm.fit", "summary.lm")]
          <- "x_lm"
nodeclass[sprof01$nodes$name == "<Anonymous>"] <- "x_Anon"
```

```
Input
sprof01$nodes$icol <-as.factor(nodeclass)
```

adds a sticky colour attribute. To interpret, you should choose your preferred colour palette, for example

ToDo: add a reference to colorbrewer

```
Input
classcol <- c("red", "green", "blue", "yellow", "cyan", "magenta", "purple")
```

ToDo: Defaults by class

ToDo: classes need separate colour palette

Nodes by package: See fig. 5 on page 18.

Nodes by class: default colour selection. See fig. 6 on page 19.

You can break down the frequency by classes of your choice. But beware of Simpson's paradox. The information you think you see may be strongly affected by your choices - what you see are reflections of conditional distributions. These may be very different from the global picture.

If package *wordcloud* is installed, a different view is possible. This is added in the plots above.

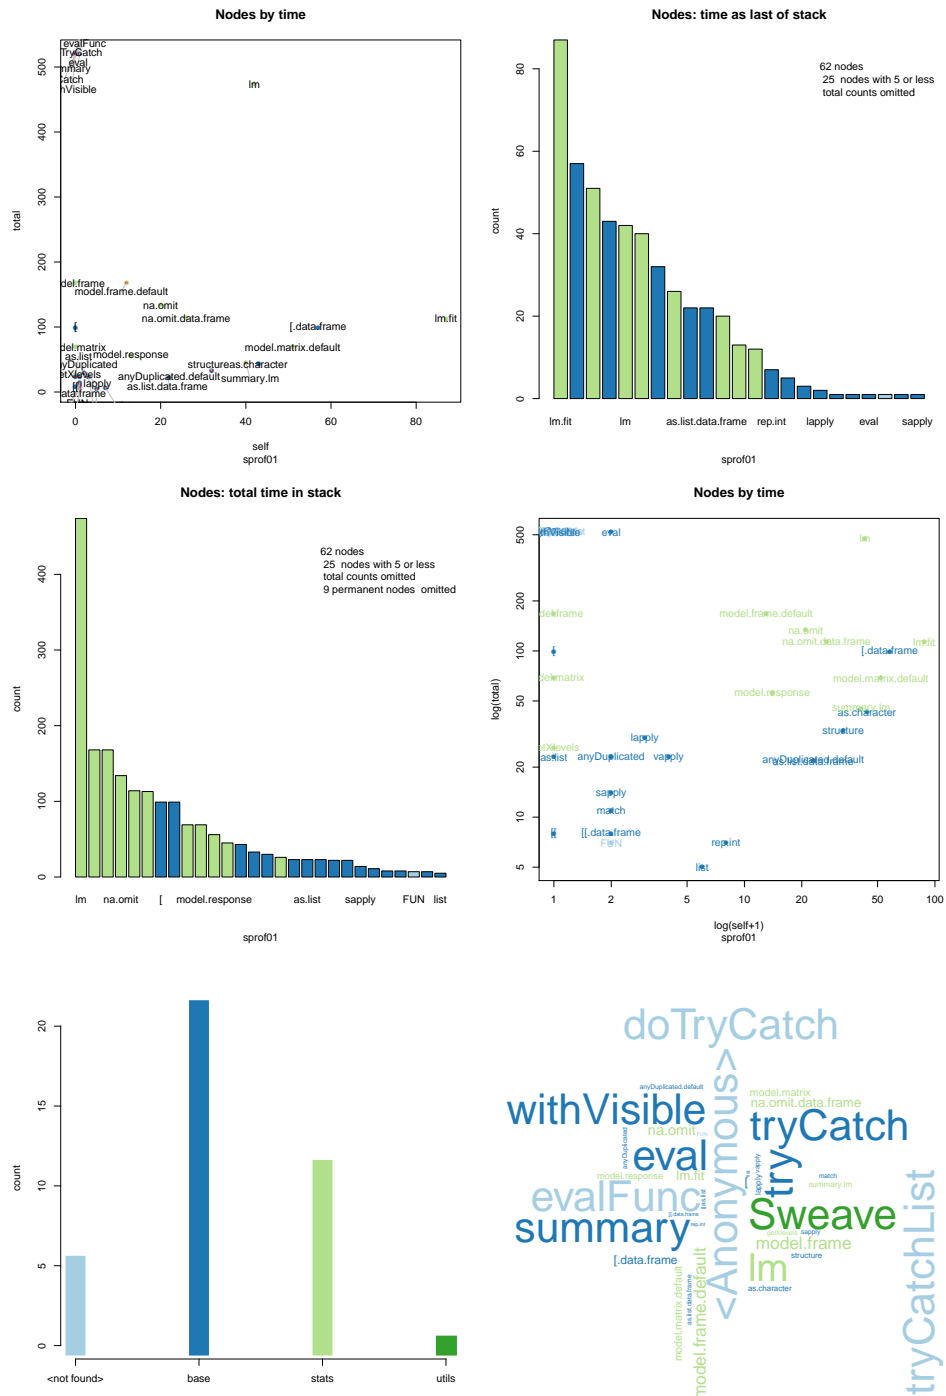
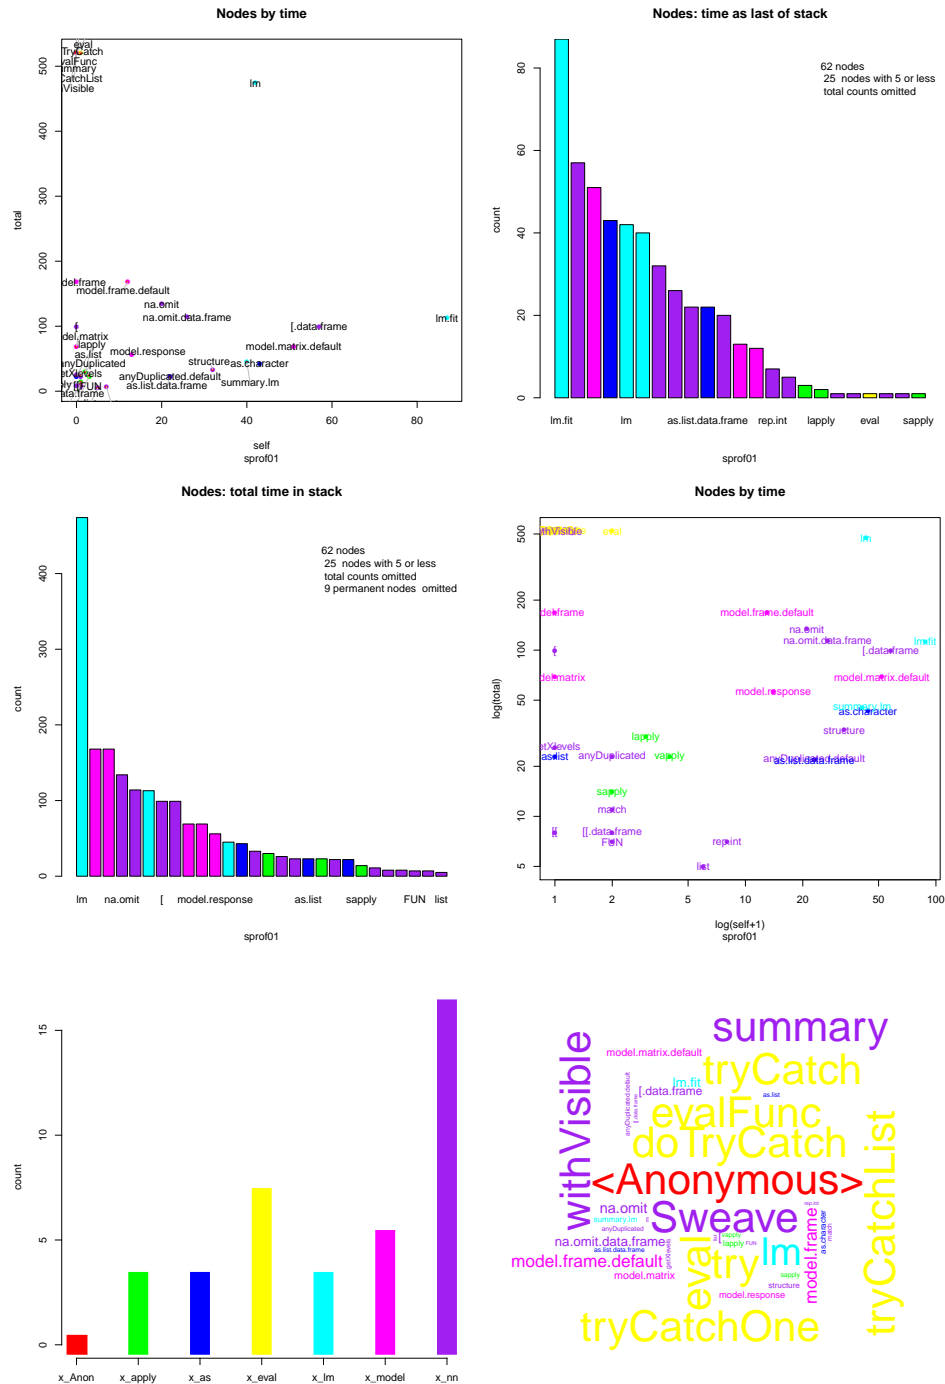


FIGURE 4. Nodes by package

```
#10 15
oldpar <- par(mfrow=c(3,2))
plot_nodes(sprof01, which=1:6, col=classcol)
par(oldpar)
```



Input

```
#10 15
oldpar <- par(mfrow=c(3,2))
plot_nodes(sprof01, which=1:6)
par(oldpar)
```

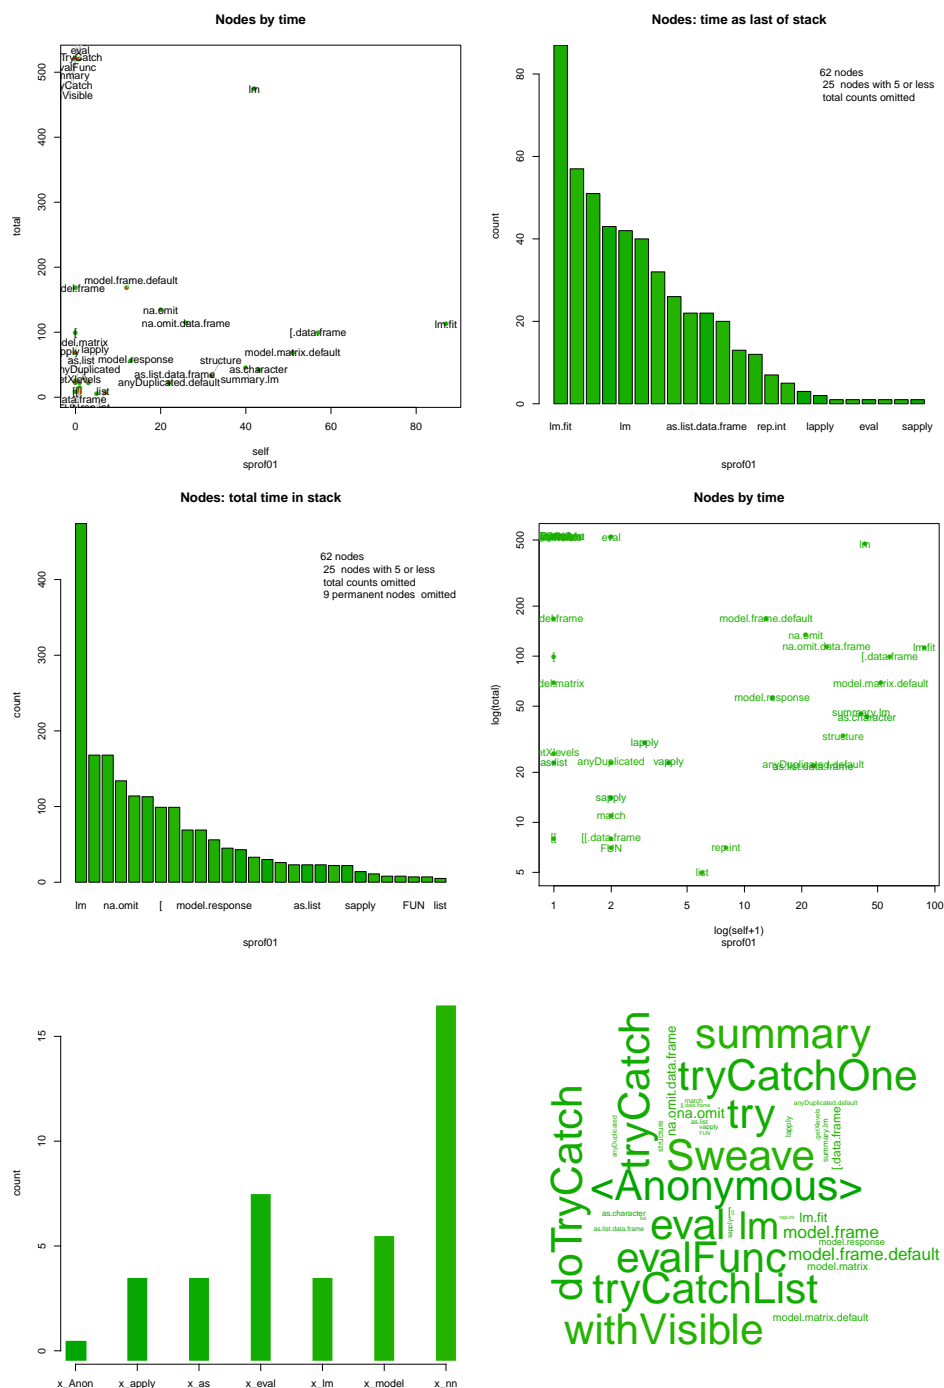


FIGURE 6. Nodes by class: default colour selection

2. A BETTER GRIP ON PROFILE INFORMATION

The basic information provided by all profilers in R is a protocol of sampled stacks. The conventional approach is to break the information down to nodes and edges. The stacks provide more information than this. One way to access it is to use linking to pass information. This has already been used on the node level in section 1.1.2 on page 9.

ToDo: add attributes to stacks, and discuss scope

ToDo: sorting/ranging stacks

2.1. The internal details. For each recorded event, the protocol records one line with a text string showing the sampled stack (in reverse order: most recent first). The stack lines may be preceded by header lines with event specific information. The protocol may be interspersed with control information, such as information about the timing interval used.

We know that the structural information, static information as well as dynamic information, can be represented with the help of a graph. For a static analysis, the graph representation may be the first choice. For a dynamic analysis, the stack information is our first information. A stack is a connected path in the program graph. If we start with nodes and edges, we lose information which is readily available in record of stacks.

As we know that we are working with stacks, we know that they have their peculiarities. Stacks tend to grow and shrink. Subsequent events will have extensions and shrinkages of stacks (if the recording is on a fine scale), or stack sharing common stumps (if the recording is on a coarser scale). We could exploit this information, but it does not seem worth the effort.

ToDo: re-think: sort stacks

There have always been interrupts, and these show up in profiles. In R, there is a related problem: garbage collection (GC) may interfere and leave traces in the stack.

Stack information is first. The call graph is a second instance that is (re)constructed from the stack recording. The graph represents cumulated one-step information. Longer scale information contained in the stacks is lost in the graph.

Here is the way we represent the profile information:

The profile log file is sanitised:

- Control lines are extracted and recorded in a separate list.
- Head parts, if present, are extracted and recorded in a matrix that is kept line-aligned with the remainder
- Line content is standardised, for example by removing stray quotation marks etc.

After this, the sanitised lines are encoded as a vector of stacks, and references to this.

If necessary, these steps are done by chunks to reduce memory load.

From the vector of stacks, a vector of nodes (or rather node names) is derived.

The stacks are now encoded by references to the nodes table. For convenience, we keep the (sanitised) textual representation of the stacks.

So far, texts are in reverse order. For each stack, we record the trailing leaf, and then we reverse order. The top of stack is now on first position.

Several statistics can be accumulated easily as a side effect.

Conceptually, the data structure consist of three tables (the implementation may differ, and is subject to change).

The profiles table is the representation of the input file. Control lines are collected in a special table. With the control lines removed, the rest is a table, one row per input line. The body of the line, the stack, is encoded as a reference to a stacks table (obligatory) and header information (optional).

The stacks table contains the collected stacks, each stack encoded as a list of references to the node table. This is obligatory. This list is kept in reverse order (root at position 1). A source line representing the stack information may be kept (optional).

The nodes table keeps the names at the nodes.

Sometimes, it is more convenient to use a simple representation, such as a matrix. Several extraction routines are provided for this, and the display routines make heavy use of this. See table 5.

ToDo: complete matrix conversion

TABLE 5. Extraction and conversion routines

<code>profiles_matrix()</code>	incidence matrix: nodes by event
<code>stacks_matrix()</code>	incidence matrix: nodes by stack
<code>list.as.matrix()</code>	fill list to equal length and convert to matrix
<code>stackstoadj()</code>	stacks to (correspondence) adjacency matrix
<code>adjacency()</code>	sprof to (correspondence) adjacency matrix

We now can go beyond node level.

This is what we get for free from the node information on our three levels: node, stack, and profile.

See fig. 7 on the next page.

ToDo: check and stabilise colour linking

The obvious message is that if seen by stack level, there are different structures. Profiling usually takes place in a framework. So at the base of the stacks, we find entries that are (almost) persistent. Then usually we have some few steps where the algorithm splits, and then we have the finer details. These can be identified using information on the stack level, but of course they are not visible on the node or edge level in a graph representation. On the stack level, we see a socket. If we want a statistic, we can look at number of different nodes by level.

Input

```
stacks_nodes <- list.as.matrix(sprof01$stacks$nodes)
nrnodes <- apply(stacks_nodes,1,function(x) {length(unique(x))})
cat("nr unique nodes per stack level\n")
```

```
#8 rainbow                                     Input
sprof01$nodes$icol <- freqranks01
shownodes(sprof01, col=rainbow(62))
```

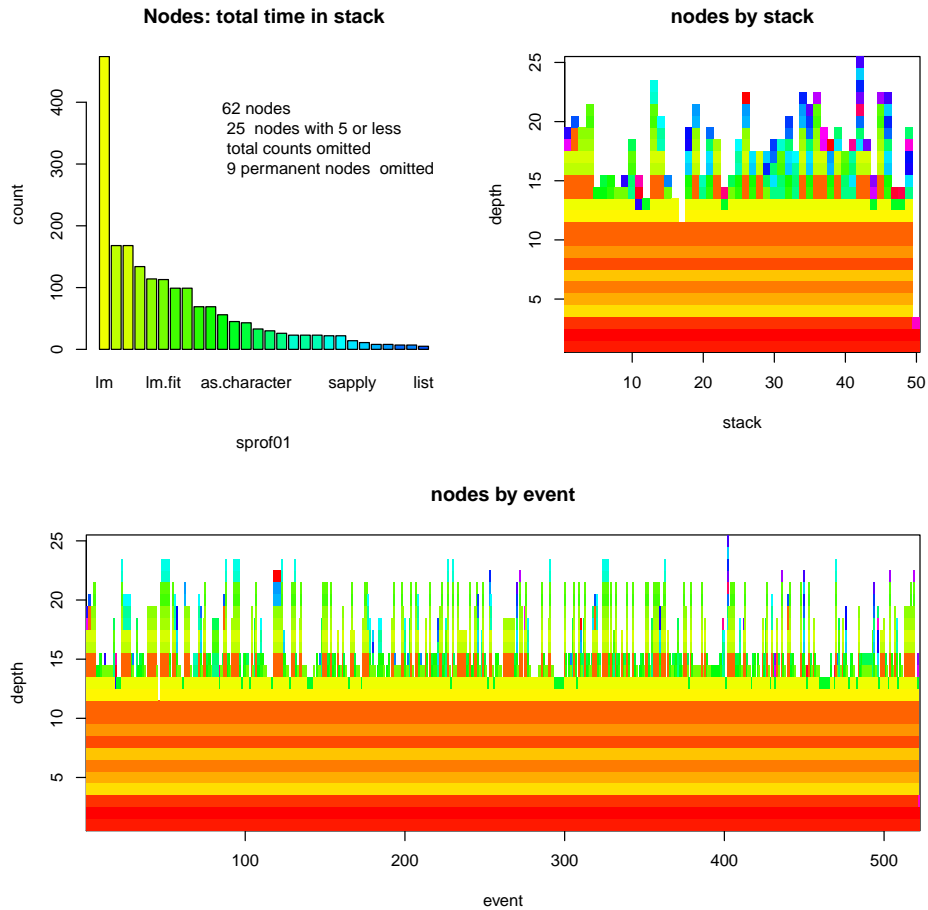


FIGURE 7. Nodes by stack and profile

nr unique nodes per stack level	Output
---------------------------------	--------

nrnodes *Input*

											Output											
[1]	1	1	2	2	2	2	2	2	2	2	2	2	4	11	12	10	10	16	9	8	6	8
[23]	3	2	2																			

 Input

```
plot(x=nrnodes, y= 1:length(nrnodes), xlab="nr of unique nodes", ylab="stack level")
abline(h=2.5,col="yellow")
abline(h=12.5,col="yellow")
```




FIGURE 8. Nr of unique nodes by stack level.

Nr. of unique nodes by stack level: See fig. 8.

We will come to finer tools in section 2.4 on page 36 but for the moment the rough information should suffice to take a decision. In our example, it is only a matter of taste whether we cut off 12 levels, or we want to work with five components after cutting 13 levels

Not so often, but a frequent phenomenon is to have some “burn in” or “fade out”. To identify this, we need to look at the profile level. The indicator to check is to whether we have very low frequency stacks at the beginning or the end of our recording. The counts to be takes as reference can be seen from the summary.

```
summary(sprof01$stacks$refcount)
```

				Output		
Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	
1.0	1.0	2.0	10.4	12.0	87.0	

ToDo: check and synchronise

```
# ----- Input -----  
  
df <- data.frame(stack=sprof01$profiles$data, count=sprof01$stacks$refcount[sprof01$profiles$data])  
#df <- xcutdata.frame(df, margin=10)  
  
prxt(df, caption="Marginal stacks", ----- Input -----  
      label="tab:margin" )
```

Table 6: Marginal stacks

	stack	count
1	1	1.00
2	2	5.00
3	2	5.00
4	3	26.00
5	4	55.00
6	4	55.00
7	5	13.00
8	5	13.00
9	6	43.00
10	7	51.00
< cut >	⋮	⋮
513	3	26.00
514	3	26.00
515	3	26.00
516	3	26.00
517	4	55.00
518	3	26.00
519	36	2.00
520	16	42.00
521	44	2.00
522	50	1.00

Here at least one recording on either side is a candidate to be off. We may have a look at the next recordings and decide to go beyond and cut off events 1 : 3 and 519 : 522.

At a closer look, we may find stack patterns (maybe marked by specific nodes) that indicate administrative intervention and rather should be handled as separators between distinct profiles rather than as part of the general dynamics. Again we may use some indicator nodes to be used as marker for special stacks.

ToDo: example

Stable framework effects sometimes are obvious and can be detected automatically. “burn in” or “fade out” may need a closer look, and special stacks need and individual

inspection on low frequency stacks. Tools for trimming are in section 2.3.1 on page 30.

2.2. The free lunch. What you have seen so far is what you get for free when using package *sprof*.


If you want to wrap up the information and look at it from a graph point of view, here is just one example. More are in section 5 on page 56. But before changing to the graph perspective, we recommend to see the next sections, not to skip them. The preview, at this point, taking package *graph* as an example:

ToDo: colours. re-colour. Propagate colour to graph.

Input

```
#6 sprofadjNEL
library(graph)
sprof01adjNEL <- as(adjacency(sprof01),"graphNEL")
plotviz(sprof01adjNEL, main="sprof01: graph layout example",
        sub=sprof01$info$id)
rm(sprof01adjNEL);
# detach("package:graph")
#! sorry. still needed by Rgraphviz -- clean up
```

sprof01: graph layout example



2.3. **Cheap thrills.** Before starting additional inspection, the data better be trimmed. Trimming routines are in section 2.3.1 on page 30, but the data structure is robust enough to allow manual intervention as used here.

ToDo: updateRprof needs careful checking. For now, we are including long listings here to provide the necessary information

Input

```
sprof02 <- sprof01; sprof02$info$id <- "sprof02: trimmed"
```

On the stack level, we take brute force to cut off the basic stacks.

Input

```
basetrim <- 13
sprof02$stacks$nodes <- sapply(sprof02$stacks$nodes,
  function (x){if (length(x)> basetrim) x[-(1:basetrim)] })
```

We have noted burn in/fade out. This is on the profile level. Taking the big knife is not advisable, since time information and stack data must be synchronised. So we are more cautious.

Input

```
sprof02$profiles$data[1:3] <- NA
sprof02$profiles$data[519:522] <- NA
```

At this point, it is a decision whether to adapt the timing information, or keep the original information. Since this decision does affect the structural information, it is not critical. But analysis is easier if unused nodes are eliminated. The *info* section is inconsistent at this point. Another reason to call `updateRprof()`.

ToDo: handle empty stacks and zero counts gracefully
ToDo: add a purge function

```
strx(sprof02$info)      Input
```

```
Output
'data.frame':      1 obs. of  8 variables:
 $ id : chr "sprof02: trimmed"
 $ date : POSIXct, format: "2013-07-30 15:17:11"
 $ nrnodes : int 62
 $ nrstacks : int 50
 $ nrrecords: int 522
 $ firstline: Factor w/ 1 level "sample.interval=1000": 1
 $ ctllines : Factor w/ 1 level "sample.interval=1000": 1
 $ ctllinenr: num 1
```

```
prxt(sprof02$nodes,cut=FALSE,
      caption="sprof02, before update",label="tab:sprof02info1")      Input
```

Table 7: sprof02, before update

	name	self.time	self.pct	total.time	total.pct	icol
1	!	2.00	0.38	2.00	0.03	50
2	..getNamespace	0.00	0.00	1.00	0.01	62
3	.deparseOpts	2.00	0.38	4.00	0.05	40
4	.getXlevels	0.00	0.00	26.00	0.34	27
5	[0.00	0.00	99.00	1.29	19
6	[.data.frame	57.00	10.92	99.00	1.29	18
7	[[0.00	0.00	8.00	0.10	35
8	[[.data.frame	1.00	0.19	8.00	0.10	36
9	%in%	1.00	0.19	4.00	0.05	41
10	<Anonymous>	6.00	1.15	522.00	6.79	1
11	\$	1.00	0.19	1.00	0.01	58
12	anyDuplicated	1.00	0.19	23.00	0.30	30
13	anyDuplicated.default	22.00	4.21	22.00	0.29	31
14	as.character	43.00	8.24	43.00	0.56	24
15	as.list	0.00	0.00	23.00	0.30	28
16	as.list.data.frame	22.00	4.21	22.00	0.29	32
17	as.list.default	1.00	0.19	1.00	0.01	57
18	as.name	1.00	0.19	1.00	0.01	56
19	coef	1.00	0.19	1.00	0.01	60
20	deparse	1.00	0.19	2.00	0.03	47
21	doTryCatch	0.00	0.00	521.00	6.78	4
22	eval	1.00	0.19	521.00	6.78	5

23	evalFunc	0.00	0.00	521.00	6.78	3
24	file	1.00	0.19	1.00	0.01	55
25	FUN	1.00	0.19	7.00	0.09	38
26	lapply	2.00	0.38	30.00	0.39	26
27	lazyLoadDBfetch	2.00	0.38	3.00	0.04	44
28	list	5.00	0.96	5.00	0.07	39
29	lm	42.00	8.05	474.00	6.16	12
30	lm.fit	87.00	16.67	113.00	1.47	17
31	match	1.00	0.19	11.00	0.14	34
32	mean	0.00	0.00	2.00	0.03	51
33	mean.default	2.00	0.38	2.00	0.03	48
34	mode	2.00	0.38	2.00	0.03	45
35	model.frame	0.00	0.00	168.00	2.18	14
36	model.frame.default	12.00	2.30	168.00	2.18	13
37	model.matrix	0.00	0.00	69.00	0.90	20
38	model.matrix.default	51.00	9.77	69.00	0.90	21
39	model.response	13.00	2.49	56.00	0.73	22
40	na.omit	20.00	3.83	134.00	1.74	15
41	na.omit.data.frame	26.00	4.98	114.00	1.48	16
42	names	2.00	0.38	2.00	0.03	49
43	NCOL	1.00	0.19	1.00	0.01	61
44	paste	0.00	0.00	1.00	0.01	59
45	pmatch	2.00	0.38	2.00	0.03	46
46	rep.int	7.00	1.34	7.00	0.09	37
47	sapply	1.00	0.19	14.00	0.18	33
48	simplify2array	0.00	0.00	4.00	0.05	43
49	structure	32.00	6.13	33.00	0.43	25
50	summary	0.00	0.00	520.00	6.76	11
51	summary.lm	40.00	7.66	45.00	0.59	23
52	Sweave	0.00	0.00	522.00	6.79	2
53	terms	0.00	0.00	2.00	0.03	52
54	terms.formula	1.00	0.19	1.00	0.01	53
55	try	0.00	0.00	521.00	6.78	10
56	tryCatch	0.00	0.00	521.00	6.78	8
57	tryCatchList	0.00	0.00	521.00	6.78	6
58	tryCatchOne	0.00	0.00	521.00	6.78	9
59	unique	3.00	0.57	4.00	0.05	42
60	unlist	0.00	0.00	1.00	0.01	54
61	vapply	3.00	0.57	23.00	0.30	29
62	withVisible	0.00	0.00	521.00	6.78	7

Input

```

sprof02 <- updateRprof(sprof02)
sprof02$info$id <- "sprof02 updated"

```

Input

```
#xtable(summary(sprof02))
#summary(sprof02)
strx(sprof02$info)
```

Output

```
'data.frame':      1 obs. of  9 variables:
 $ id : chr "sprof02 updated"
 $ date : POSIXct, format: "2013-07-30 15:17:11"
 $ nrnodes : int 62
 $ nrstacks : int 50
 $ nrrecords : int 522
 $ firstline : Factor w/ 1 level "sample.interval=1000": 1
 $ ctllines : Factor w/ 1 level "sample.interval=1000": 1
 $ ctllinenr : num 1
 $ date_updated: POSIXct, format: "2013-07-30 15:17:13"
```

Input

```
prxt(sprof02$nodes, cut=FALSE, caption="sprof02, after update",
     label="tab:sprof02info2")
```

Table 8: sprof02, after update

	name	self.time	self.pct	total.time	total.pct	icol
1	!	1.00	0.23	1.00	0.06	50
2	..getNamespace	0.00	0.00	1.00	0.06	62
3	.deparseOpts	2.00	0.46	4.00	0.25	40
4	.getXlevels	0.00	0.00	26.00	1.64	27
5	[0.00	0.00	98.00	6.17	19
6	[.data.frame	57.00	13.16	98.00	6.17	18
7	[[0.00	0.00	8.00	0.50	35
8	[[.data.frame	1.00	0.23	8.00	0.50	36
9	%in%	1.00	0.23	4.00	0.25	41
10	<Anonymous>	6.00	1.39	6.00	0.38	1
11	\$	1.00	0.23	1.00	0.06	58
12	anyDuplicated	1.00	0.23	23.00	1.45	30
13	anyDuplicated.default	22.00	5.08	22.00	1.39	31
14	as.character	43.00	9.93	43.00	2.71	24
15	as.list	0.00	0.00	23.00	1.45	28
16	as.list.data.frame	22.00	5.08	22.00	1.39	32
17	as.list.default	1.00	0.23	1.00	0.06	57
18	as.name	1.00	0.23	1.00	0.06	56
19	coef	1.00	0.23	1.00	0.06	60
20	deparse	1.00	0.23	2.00	0.13	47
21	doTryCatch	0.00	0.00	0.00	0.00	4
22	eval	0.00	0.00	164.00	10.33	5
23	evalFunc	0.00	0.00	0.00	0.00	3
24	file	0.00	0.00	0.00	0.00	55
25	FUN	1.00	0.23	7.00	0.44	38
26	lapply	2.00	0.46	30.00	1.89	26
27	lazyLoadDBfetch	1.00	0.23	1.00	0.06	44

28	list	3.00	0.69	3.00	0.19	39
29	lm	0.00	0.00	0.00	0.00	12
30	lm.fit	87.00	20.09	113.00	7.12	17
31	match	1.00	0.23	11.00	0.69	34
32	mean	0.00	0.00	1.00	0.06	51
33	mean.default	1.00	0.23	1.00	0.06	48
34	mode	2.00	0.46	2.00	0.13	45
35	model.frame	0.00	0.00	164.00	10.33	14
36	model.frame.default	12.00	2.77	164.00	10.33	13
37	model.matrix	0.00	0.00	69.00	4.35	20
38	model.matrix.default	51.00	11.78	69.00	4.35	21
39	model.response	13.00	3.00	56.00	3.53	22
40	na.omit	20.00	4.62	133.00	8.38	15
41	na.omit.data.frame	26.00	6.00	113.00	7.12	16
42	names	2.00	0.46	2.00	0.13	49
43	NCOL	1.00	0.23	1.00	0.06	61
44	paste	0.00	0.00	1.00	0.06	59
45	pmatch	2.00	0.46	2.00	0.13	46
46	rep.int	7.00	1.62	7.00	0.44	37
47	sapply	1.00	0.23	14.00	0.88	33
48	simplify2array	0.00	0.00	4.00	0.25	43
49	structure	32.00	7.39	33.00	2.08	25
50	summary	0.00	0.00	0.00	0.00	11
51	summary.lm	0.00	0.00	0.00	0.00	23
52	Sweave	0.00	0.00	0.00	0.00	2
53	terms	0.00	0.00	1.00	0.06	52
54	terms.formula	1.00	0.23	1.00	0.06	53
55	try	0.00	0.00	0.00	0.00	10
56	tryCatch	0.00	0.00	0.00	0.00	8
57	tryCatchList	0.00	0.00	0.00	0.00	6
58	tryCatchOne	0.00	0.00	0.00	0.00	9
59	unique	3.00	0.69	4.00	0.25	42
60	unlist	0.00	0.00	1.00	0.06	54
61	vapply	3.00	0.69	23.00	1.45	29
62	withVisible	0.00	0.00	0.00	0.00	7

Input

See fig. 9 on the next page.

ToDo: This section
needs to be reworked

ToDo: trimexample

ToDo: add trim by
keyword

2.3.1. *Trimming.*

Input

```
trimstacks <- function(sprof, level){
  lapply(sprof$stacks$nodes, function(x) {x[-(1:level)]})
}
```

#8 *Input*

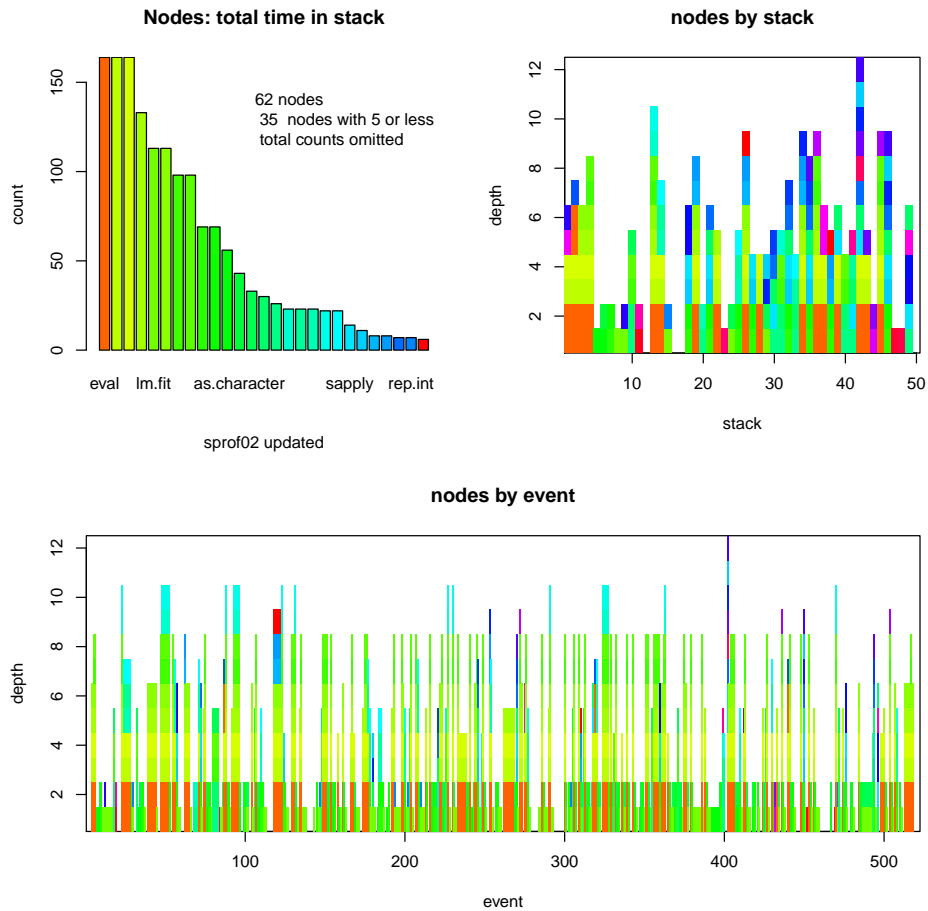
`shownodes(sprof02)`


FIGURE 9. Nodes by stack and profile

Input

```
sprof01Tr <- trimstacks(sprof01, 11)
#profile_nodesTr <- profiles_matrix(sprof01Tr)
#image(x=1:ncol(profile_nodesTr),y=1:nrow(profile_nodesTr), t(profile_nodesTr),xlab="event", ylab="
```


Input

```
nodefreq <- rep(0,length(sprof01$nodes$name))
for (i in (1:length(sprof01$stacks$nodes))){
  nodefreq <- nodefreq +
    table( factor(sprof01$stacks$nodes[[i]],
                  levels <- 1:length(sprof01$nodes$name),
                  ordered=FALSE))
```


```
}
names(nodefreq) <- sprof01$nodes$name
```

Top frequent nodes.

```
ndf <- nodefreq[nodefreq>1]
ondf <- order(ndf,decreasing=TRUE)
barplot(ndf[ondf])
```



```
barplot(ndf[ondf], col=rainbow(length(ondf)))
```

Top frequent stacks.

Input

```

x <- sprof01
xsrc <- as.matrix(x$stacks$refcount)
rownames(xsrc) <- rownames(xsrc, do.NULL=FALSE, prefix="S")
#stf <- x$stacks$refcount[x$stacks$refcount>1]
#names(stf) <- x$stacks$shortname[x$stacks$refcount>1]
stf <- xsrc[xsrc>1]
names(stf) <- rownames(xsrc)[xsrc>1]
ostf <- order(stf, decreasing=TRUE)
barplot(stf[ostf])

```

ToDo: cut next level

Control structures may be represented in R as function, and these may lead to concentration points. Using information from the stacks, we can avoid these by introducing substitute nodes on the stack level. For example, `lapply` is appearing in various contexts and may be confusing any graph representation. We can avoid this by replacing a short sequence. `"[" "lapply" ".getXlevels" -> "<.getXlevels_[">` If the node does not exist, we want to add it to our global variable. For now, we do it using expressions on the R basic level and avoid tricks like simulating “call by reference”.

ToDo: Implement. Currently best handled on source=text level
ToDo: function addnode using “call by reference” to be added

Input

```

sprof03 <- sprof02; sprof03$info$id <- "sprof03: surgery"
node <- "<.getXlevels_[">
#nodei <- function(sprofx, node, warn = TRUE)
{
  i <- match(node, sprof03$nodes$name, nomatch=0)
  if (i==0){
    sprof03$nodes$name <- as.character(sprof03$nodes$name)
    sprof03$nodes <- rbind(sprof03$nodes,NA)
    i <- length(sprof03$nodes$name)
    sprof03$nodes$name[i] <- node
    if (as.logical(options("warn")))
      message("node added. An updateRprof() may be necessary.")
  }
  nodei <- i
}
# sprof <- sprof01; nodei(sprof,"kiki"); sprof$nodes

```

Now we have to identify the stacks that may get a replacement. First find the candidates.

Input

```

targeti <- match("lapply", sprof03$nodes$name, nomatch=0)
found <- lapply(sprof03$stacks$nodes, function(X) match(targeti,X))
found <- data.frame(stack=1:length(found), position=as.matrix(found))
found[!is.na(found$position),]

```

Output

	stack	position
10	10	2
21	21	5
30	30	3
31	31	2
32	32	5
33	33	2
35	35	5
39	39	6
40	40	2
41	41	3
42	42	6
46	46	5
49	49	6

Input

```
# as.factor(sprof03$stacks$nodes[!is.na(found)], levels=1: length(sprof03$nodes$name), labels=sprof03$nodes$name)
```

ToDo: implement replacement on the stack level.

ToDo: implement

For now, these are just candidates.

Other candidates are: "as.list" "vapply" "model.frame.default" -> "<model_as.list>" or "as.list" "vapply" "model.matrix.default" -> "<model_matrix_as.list>"

```

newchopnode <- function(nodenames, chop) {
  tmpname <- paste("<", as.character(nodenames[chop]), ">")
  # chec for existing.
  # add if necessary
  tmpname
}
chopstack <- function(x, chop, replacement)
{
  # is chop in x`
  # y: cut x.
  # merge x <- head + replacement + tail
  return(x)
}

```

ToDo: needs serious revision

ToDo: support na removal in rrle

2.4. Run length. For a visual inspection, runs of the same node and level in the profile are easily perceived. For an analytical inspection, we have to reconstruct the runs from the data. In stacks, runs are organised hierarchically. On the root level, runs are just ordinary runs. On the next levels, runs have to be defined given (within) the previous runs. So we need a recursive version of **rrle**, applied to the profile information. This gives a detailed information about the presence time of each node, by stack level.

ToDo: use sprof02 or sprof03?

```

profile_nodes <- profiles_matrix(sprof02)
profile_nodes_rrle <- rrle(profile_nodes)
strx(profile_nodes_rrle)

```

```

List of 12
$ :List of 2
..$ lengths: int [1:365] 1 1 1 3 3 1 7 1 ...
..$ values : int [1:365] NA NA NA 22 39 37 30 4 ...
..- attr(*, "class")= chr "rrle"
$ :List of 2
..$ lengths: int [1:411] 1 1 1 3 1 1 1 1 ...
..$ values : int [1:411] NA NA NA 22 NA NA 14 38 ...
..- attr(*, "class")= chr "rrle"
$ :List of 2
..$ lengths: int [1:431] 1 1 1 3 1 1 1 1 ...
..$ values : int [1:431] NA NA NA 35 NA NA NA NA ...
..- attr(*, "class")= chr "rrle"
$ :List of 2
..$ lengths: int [1:431] 1 1 1 3 1 1 1 1 ...
..$ values : int [1:431] NA NA NA 36 NA NA NA NA ...

```

```

..- attr(*, "class")= chr "rle"
$ :List of 2
..$ lengths: int [1:452] 1 1 1 3 1 1 1 1 ...
..$ values : int [1:452] NA NA NA 40 NA NA NA NA ...
..- attr(*, "class")= chr "rle"
$ :List of 2
..$ lengths: int [1:468] 1 1 1 3 1 1 1 1 ...
..$ values : int [1:468] NA NA NA 41 NA NA NA NA ...
..- attr(*, "class")= chr "rle"
$ :List of 2
..$ lengths: int [1:490] 1 1 1 1 2 1 1 1 ...
..$ values : int [1:490] NA NA NA NA 5 NA NA NA ...
..- attr(*, "class")= chr "rle"
$ :List of 2
..$ lengths: int [1:494] 1 1 1 1 2 1 1 1 ...
..$ values : int [1:494] NA NA NA NA 6 NA NA NA ...
..- attr(*, "class")= chr "rle"
$ :List of 2
..$ lengths: int [1:508] 1 1 1 1 1 1 1 1 ...
..$ values : int [1:508] NA NA NA NA NA NA NA NA ...
..- attr(*, "class")= chr "rle"
$ :List of 2
..$ lengths: int [1:512] 1 1 1 1 1 1 1 1 ...
..$ values : int [1:512] NA NA NA NA NA NA NA NA ...
..- attr(*, "class")= chr "rle"
$ :List of 2
..$ lengths: int [1:522] 1 1 1 1 1 1 1 1 ...
..$ values : int [1:522] NA NA NA NA NA NA NA NA ...
..- attr(*, "class")= chr "rle"
$ :List of 2
..$ lengths: int [1:522] 1 1 1 1 1 1 1 1 ...
..$ values : int [1:522] NA NA NA NA NA NA NA NA ...
..- attr(*, "class")= chr "rle"

```

Input

```

collapseNA0 <- function(x)
{
  xv <- is.na(x$values)
  xve <- rle(xv)
  xve$values[xve$lengths==1] <- FALSE # ignore single NAs
  xve$pos <- cumsum(xve$lengths) # pos is end position
  xve
  ### start replacement here
}
c0 <- collapseNA0(profile_nodes_rle[[10]])
str(c0)

```

Output

```

List of 3
 $ lengths: int [1:27] 22 1 24 1 35 1 4 1 26 1 ...
 $ values : logi [1:27] TRUE FALSE TRUE FALSE TRUE FALSE ...
 $ pos    : int [1:27] 22 23 47 48 83 84 88 89 115 116 ...
- attr(*, "class")= chr "rle"

```

On a given stack level, the run length is the best information on the time used per call, and the run count of a node is the best information on the number of calls. So this is a prime starting point for in-depth analysis.

ToDo: keep as factor. This is a sparse cube with margins node, stack level, run length. Nodes are mostly concentrated on few levels.
ToDo: Warning: data structure still under discussion

Input

```
# side effect: NAs are removed
profile_nodes_rlet <- lapply(profile_nodes_rle,
  function(x) table(x,dnn=c("run length","node"))) )
invisible(lapply(profile_nodes_rlet,
  function(x) print.table(x,zero.print = ".") ))
```

Output

```
      node
run length  2  4 11 19 22 30 32 37 39 43
  1  1 17  1  1 40 46  1 55 35  1
  2  .  1  .  . 17 18  .  4  3  .
  3  .  1  .  .  7  3  .  2  3  .
  4  .  1  .  .  4  1  .  .  .  .
  5  .  .  .  .  2  1  .  .  .  .
  6  .  .  .  .  6  1  .  .  1  .
  7  .  .  .  .  1  1  .  .  .  .

      node
run length 14 18 22 26 27 33 38 46 47 49 61
  1 34  1 40 16  1  1 55  7  3 10  .
  2  3  . 17  .  .  .  4  .  .  1  .
  3  1  .  7  .  .  .  2  .  .  .  1
  4  .  .  4  1  .  .  .  .  .  .  .
  5  .  .  2  .  .  .  .  .  .  .  .
  6  .  .  6  .  .  .  .  .  .  1  .
  7  .  .  1  .  .  .  .  .  .  .  .

      node
run length  5  7  9 15 26 31 35 48 61
  1 14  1  1  2  2  9 40  1  6
  2  .  .  .  .  .  . 17  .  1
  3  .  .  .  .  .  .  7  .  .
  4  1  .  .  .  .  .  4  .  .
  5  .  .  .  .  .  .  2  .  .
  6  .  .  .  .  .  .  6  .  .
  7  .  .  .  .  .  .  1  .  .

      node
run length  6  8 15 16 25 31 36 47 59
  1 14  1  7  2  1  1 40  9  1
  2  .  .  1  .  .  . 17  .  .
  3  .  .  .  .  .  .  7  .  .
  4  1  .  .  .  .  .  4  .  .
  5  .  .  .  .  .  .  2  .  .
  6  .  .  .  .  .  .  6  .  .
  7  .  .  .  .  .  .  1  .  .

      node
run length  3 10 12 16 17 20 22 26 40 47 48 49 53 60 61
  1  1  1  1  6  1  1  3  5 46  2  3 11  1  1  7
  2  .  .  .  1  .  .  .  . 10  .  .  .  .  .  .
  3  .  .  .  .  .  .  .  .  5  .  .  .  .  .  .
```

```

      4 . . . . . . . . . 4 . . 1 . . .
      5 . . . . . . . . . 1 . . . . . 1
      6 . . . . . . . . . 4 . . . . . .
      7 . . . . . . . . . 1 . . . . . .
      node
run length 15 22 25 26 41 54 59
      1 7 3 5 3 43 1 3
      2 . . . . 7 . .
      3 . . . . 4 . .
      4 . . . . 5 . .
      5 1 . . . . . .
      6 . . . . 4 . .
      node
run length 3 5 7 9 16 25 28
      1 3 45 2 1 7 1 3
      2 . 4 . . . . .
      3 . 3 . . . . .
      4 . 3 . . . . .
      5 . . 1 . 1 . .
      6 . 1 . . . . .
      node
run length 6 8 31 44 45
      1 45 2 1 1 2
      2 4 . . . . .
      3 3 . . . . .
      4 3 . . . . .
      5 . 1 . . . .
      6 1 . . . . .
      node
run length 1 9 10 12 20 34 42
      1 1 1 . 9 1 1 2
      4 . . . 2 . . .
      5 . . 1 1 . . .
      node
run length 9 13
      1 1 9
      4 . 2
      5 . 1
      node
run length 31
      1 1
      node
run length 34
      1 1

```

This is a poor first attempt to tame *profile_nodes_rlet*.

Input

```

maxnode <-0
maxlen <-0
maxlevel <-length(profile_nodes_rle)
for (lev in (1:maxlevel) ) {
  proflev <- profile_nodes_rle[[lev]]

```

ToDo: some mess in the code below – not working

ToDo: replace by decent vector/array based implementation

ToDo: add names for node dimension

```

    if (!is.null(proflev)) {
      maxn <- max(proflev$values, na.rm=TRUE)
      if (maxn>maxnode) maxnode <- maxn
      maxl <- max(proflev$lengths, na.rm=TRUE)
      if (maxl>maxlen) maxlen <- maxl
      # cat("Level ",lev,maxn," Length:",maxl,"\n")
    }
  }
  ## collapse profile_nodes_rle to 3d array. Allocate memory first.
  profile_nodes_rlearray <- array(0,
    dim=c(maxnode,length(profile_nodes_rle), maxlen),
    dimnames= list("node"=sprof02$nodes$name[1:maxnode],
      "level"=1:length(profile_nodes_rle),
      "run_length"=1:maxlen))
  strx(profile_nodes_rlearray)

```

Output

```

num [1:61, 1:12, 1:7] 0 0 0 0 0 0 0 0 0 ...
- attr(*, "dimnames")=List of 3
..$ node : chr [1:61] "!" "..getNamespace" ".deparseOpts" ...
..$ level : chr [1:12] "1" "2" "3" ...
..$ run_length: chr [1:7] "1" "2" "3" ...

```

Input

```

for (lev in (1:maxlevel) ) {
  proflev <- profile_nodes_rle[[lev]]
  if (!is.null(proflev)) {
    for (j in (1: length(proflev$lengths))) {
      if (!is.na(proflev$values[j])) {
        profile_nodes_rlearray[proflev$values[j],lev,proflev$lengths[j]] <-
        profile_nodes_rlearray[proflev$values[j],lev,proflev$lengths[j]] +1
        #cat(lev,j,":",proflev$values[j],lev,proflev$lengths[j],"\n")
      }#if (!is.na
    }#for j
  }
}

```

ToDo: add
marginals and
conditionals.
Provide function
node_summary.

This allows us to extract marginal from `provlev[node, level, run length]`.

Input

```

node=41
cat(as.character(sprof02$nodes$name[node]),"\n")

```

Output

```

na.omit.data.frame

```

Input

```

nn <- profile_nodes_rlearray[node, , ]
#dimnames(nn)<- list( "level", "run_length")
#nn
print.table(addmargins(nn), zero.print = ".")

```


Output							
	run_length						
level	1	2	3	4	5	6	7 Sum
1
2
3
4
5
6	43	7	4	5	.	4	63
7
8
9
10
11
12
Sum	43	7	4	5	.	4	63

```

Input
mt <- margin.table(profile_nodes_rlearray, margin = c(1,3))
amt <- addmargins(mt)
amt <- amt[amt[, "Sum"]>0,]
print.table(amt, zero.print = ".")

```

ToDo: rescale to application scale

ToDo: replace sum by weighted sum

ToDo: allow sorting, e.g. by marginals

Output							
	run_length						
node	1	2	3	4	5	6	7 Sum
!	1	1
..getNamespace	1	1
.deparseOpts	4	4
.getXlevels	17	1	1	1	.	.	20
[59	4	3	4	.	1	71
[.data.frame	59	4	3	4	.	1	71
[[3	.	.	.	1	.	4
[[.data.frame	3	.	.	.	1	.	4
%in%	4	4
<Anonymous>	1	.	.	.	1	.	2
\$	1	1
anyDuplicated	10	.	.	2	1	.	13
anyDuplicated.default	9	.	.	2	1	.	12
as.character	34	3	1	.	.	.	38
as.list	16	1	.	.	1	.	18
as.list.data.frame	15	1	.	.	1	.	17
as.list.default	1	1
as.name	1	1
coef	1	1
deparse	2	2
eval	86	34	14	8	4	12	160
FUN	7	7
lapply	26	.	.	1	.	.	27
lazyLoadDBfetch	1	1
list	3	3
lm.fit	46	18	3	1	1	1	71
match	12	12

mean	1	1
mean.default	1	1
mode	2	2
model.frame	40	17	7	4	2	6	1	77
model.frame.default	40	17	7	4	2	6	1	77
model.matrix	55	4	2	61
model.matrix.default	55	4	2	61
model.response	35	3	3	.	.	1	.	42
na.omit	46	10	5	4	1	4	1	71
na.omit.data.frame	43	7	4	5	.	4	.	63
names	2	2
NCOL	1	1
paste	1	1
pmatch	2	2
rep.int	7	7
sapply	14	14
simplify2array	4	4
structure	21	1	.	1	.	1	.	24
terms	1	1
terms.formula	1	1
unique	4	4
unlist	1	1
vapply	13	1	1	.	1	.	.	16
Sum	813	130	56	41	18	37	6	1101

These are some attempts to recover the factor structures.

```

Input
xfi <- levels(sprof02$nodes$name)
profile_nodes_rlefac <- lapply(profile_nodes_rle,
  function(xl) {xl$values <- factor(xl$values,
    levels=1:62,
    labels=xfi); xl}) # seems ok
profile_nodes_rletfac <- lapply(profile_nodes_rle,
  function(x) table(x,dnn=c("run length","node")) ) #factors lost again
colnames(profile_nodes_rletfac[[1]]) <-
  sprof02$nodes$name[ as.integer(colnames(profile_nodes_rletfac[[1]]))]
profile_nodes_rletfac1 <- lapply(profile_nodes_rletfac,
  function(xl) {colnames(xl) <-
    sprof02$nodes$name[ as.integer(colnames(xl))];
    xl} )
invisible(lapply(profile_nodes_rletfac1,
  function(x) print.table(t(x),zero.print = ".") ))

```

```

Output
      run length
node  1  2  3  4  5  6  7
<NA>  1  .  .  .  .  .  .
<NA> 17  1  1  1  .  .  .
<NA>  1  .  .  .  .  .  .
<NA>  1  .  .  .  .  .  .
<NA> 40 17  7  4  2  6  1
<NA> 46 18  3  1  1  1  1

```

```

<NA> 1 . . . . .
<NA> 55 4 2 . . . .
<NA> 35 3 3 . . 1 .
<NA> 1 . . . . .

run length
node      1 2 3 4 5 6 7
as.character 34 3 1 . . . .
as.name      1 . . . . .
eval        40 17 7 4 2 6 1
lapply      16 . . 1 . . .
lazyLoadDBfetch 1 . . . . .
mean.default 1 . . . . .
model.matrix.default 55 4 2 . . . .
rep.int      7 . . . . .
sapply       3 . . . . .
structure    10 1 . . . 1 .
vapply       . . 1 . . . .

run length
node      1 2 3 4 5 6 7
[         14 . . 1 . . .
[[        1 . . . . .
%in%      1 . . . . .
as.list   2 . . . . .
lapply    2 . . . . .
match     9 . . . . .
model.frame 40 17 7 4 2 6 1
simplify2array 1 . . . . .
vapply    6 1 . . . . .

run length
node      1 2 3 4 5 6 7
[.data.frame 14 . . 1 . . .
[ [.data.frame 1 . . . . .
as.list       7 1 . . . . .
as.list.data.frame 2 . . . . .
FUN           1 . . . . .
match         1 . . . . .
model.frame.default 40 17 7 4 2 6 1
sapply        9 . . . . .
unique        1 . . . . .

run length
node      1 2 3 4 5 6 7
.deparseOpts 1 . . . . .
<Anonymous>  1 . . . . .
anyDuplicated 1 . . . . .
as.list.data.frame 6 1 . . . . .
as.list.default 1 . . . . .
deparse       1 . . . . .
eval          3 . . . . .
lapply        5 . . . . .
na.omit       46 10 5 4 1 4 1
sapply        2 . . . . .
simplify2array 3 . . . . .
structure     11 . . 1 . . .

```

```

terms          1 . . . . .
unlist         1 . . . . .
vapply         7 . . . 1 . .
               run length
node           1 2 3 4 5 6
as.list        7 . . . 1 .
eval           3 . . . . .
FUN            5 . . . . .
lapply         3 . . . . .
na.omit.data.frame 43 7 4 5 . 4
terms.formula  1 . . . . .
unique         3 . . . . .
               run length
node           1 2 3 4 5 6
.deparseOpts   3 . . . . .
[              45 4 3 3 . 1
[[             2 . . . 1 .
%in%           1 . . . . .
as.list.data.frame 7 . . . 1 .
FUN            1 . . . . .
list           3 . . . . .
               run length
node           1 2 3 4 5 6
[.data.frame  45 4 3 3 . 1
[[.data.frame  2 . . . 1 .
match          1 . . . . .
paste          1 . . . . .
pmatch         2 . . . . .
               run length
node           1 4 5
!              1 . .
%in%           1 . .
<Anonymous>    . . 1
anyDuplicated  9 2 1
deparse        1 . .
mode           1 . .
names          2 . .
               run length
node           1 4 5
%in%           1 . .
anyDuplicated.default 9 2 1
               run length
node           1
match 1
               run length
node           1
mode 1

```

Input

ToDo: add current level

ToDo: generate a coplot representation

ToDo: add time per call information: add marginals statistics run time by node

ToDo: table: node #runs min median run length max

3. XXX – LOST & FOUND

```

Input
# xtable cannot handle posix, so we use print output here
# str(profile_nodes_rle, max.level=2, vec.len=3, nchar.max=40, list.len=6)
strx(sprof01$info)

```

```

Output
'data.frame':      1 obs. of  8 variables:
 $ id : chr "sprof01"
 $ date : POSIXct, format: "2013-07-30 15:17:11"
 $ nrnodes : int 62
 $ nrstacks : int 50
 $ nrrecords: int 522
 $ firstline: Factor w/ 1 level "sample.interval=1000": 1
 $ ctllines : Factor w/ 1 level "sample.interval=1000": 1
 $ ctllinenr: num 1

```


Selections are recorded as selection vectors, with reference to the original order. This needs some caution to align them with the order choices.

```

Input
rownames(sprof01$nodes) <- sprof01$nodes$names
nodesperm <- order(sprof01$nodes$total.time,decreasing=TRUE)
nodesnrobsok <- sprof01$nodes$total.time > 4
sp <- sprof01$nodes$total.time[nodesperm][nodesnrobsok[nodesperm]]
names(sp) <- sprof01$nodes$name[nodesperm][nodesnrobsok[nodesperm]]
barplot(sp,
  main="Nodes, by total time", ylab="total time")

```

Nodes, by total time




On the first look, information on the profile level is not informative. Profile records are just recordings of some step, taken at regular intervals. We get a minimal information, if we encode the stacks in colour.

ToDo: use stack colours

We now do a step down analysis. Aggregating the information from the profiling events, we have the frequency of stack references. On the stack level, we encode the frequency in colour, and linking propagates this to the profile level.

Input

```
stackfreqscore <- rank(sprof01$stacks$refcount,ties.method="random")
stacksperm <- order(sprof01$stacks$refcount,decreasing=TRUE)
stacksnrobsok <- sprof01$stacks$refcount > 4
stackfreqscore4<- stackfreqscore[stacksperm][stacksnrobsok[stacksperm]]
barplot(sp[stacksnrobsok[stacksperm]],
        main="Stacks, by reference count (4 obs. minimum)", ylab="count > 4",
        col=rainbow(80)[stackfreqscore4], sub=sprof01$info$id)
```



Input

```
prxt(sprof01$nodes,
     caption="nodes",
     label="tab:prSREnodes",
     max.level=2, vec.len=3,nchar.max=40)
```

Table 9: nodes

	name	self.time	self.pct	total.time	total.pct	icol
1	!	2.00	0.38	2.00	0.03	50
2	..getNamespace	0.00	0.00	1.00	0.01	62
3	.deparseOpts	2.00	0.38	4.00	0.05	40
4	.getXlevels	0.00	0.00	26.00	0.34	27
5	[0.00	0.00	99.00	1.29	19
6	[.data.frame	57.00	10.92	99.00	1.29	18
7	[[0.00	0.00	8.00	0.10	35
8	[[.data.frame	1.00	0.19	8.00	0.10	36
9	%in%	1.00	0.19	4.00	0.05	41
10	<Anonymous>	6.00	1.15	522.00	6.79	1
< cut >	\vdots	⋮	⋮	⋮	⋮	⋮
53	terms	0.00	0.00	2.00	0.03	52
54	terms.formula	1.00	0.19	1.00	0.01	53
55	try	0.00	0.00	521.00	6.78	10
56	tryCatch	0.00	0.00	521.00	6.78	8
57	tryCatchList	0.00	0.00	521.00	6.78	6
58	tryCatchOne	0.00	0.00	521.00	6.78	9
59	unique	3.00	0.57	4.00	0.05	42
60	unlist	0.00	0.00	1.00	0.01	54
61	vapply	3.00	0.57	23.00	0.30	29
62	withVisible	0.00	0.00	521.00	6.78	7

Input

```
#str(sprof01$stacks, max.level=2, vec.len=6,
#      nchar.max=40, list.len=20,width=70, strict.width="wrap"
strx(sprof01$stacks)
```

Output

```
'data.frame':      50 obs. of  7 variables:
 $ nodes :List of 50
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 ..$ : int 52 10 23 55 56 57 58 21 ...
 .. [list output truncated]
 $ shortname : Factor w/ 50 levels
   "S<A>eFttCtCLtCOdTCwVeesleem.m..n.n...["| __truncated__,...: 27 17
   19 1 35 36 37 30 ...
 $ refcount : num 1 5 26 55 13 43 51 87 ...
 $ stacklength : int 19 20 19 21 14 15 15 14 ...
 $ stackheadnodes: int 52 52 52 52 52 52 52 52 ...
 $ stackleafnodes: int 27 28 41 6 39 14 38 30 ...
 $ stackssrc : Factor w/ 50 levels "!" [.data.frame [
   na.omit.data.frame na."| __truncated__,...: 27 28 39 5 37 13 36 30
   ...
```

A summary is provided on request.

Input

```
sumsprof01 <- summary.sprof(sprof01)
```

Output

```
$id
[1] "Profile Summary Tue Jul 30 15:17:13 2013"

$len
[1] 522

$uniquestacks
[1] 50

$nr_runs
[1] 396

$nrstacks
[1] 50

$stacklength
```

[1] 3 25

\$nrnodesperlevel

[1] 1 1 2 1 1 1 1 1 1 1 1 1 3 10 11 9 9 15 8 7 5 7
 [23] 2 1 1

	shortname	root	leaf	self.time	self.pct
!	!	-	LEAF	2	0.383142
..getNamespace	..gN	-	-	0	0.000000
.deparseOpts	.dpO	-	LEAF	2	0.383142
.getXlevels	.gtX	-	-	0	0.000000
[[-	-	0	0.000000
[.data.frame	[.d.	-	LEAF	57	10.919540
[[[[-	-	0	0.000000
[[.data.frame	[[..	-	LEAF	1	0.191571
%in%	%in%	-	LEAF	1	0.191571
<Anonymous>	<An>	-	LEAF	6	1.149425
\$	\$	-	LEAF	1	0.191571
anyDuplicated	anyD	-	LEAF	1	0.191571
anyDuplicated.default	anD.	-	LEAF	22	4.214559
as.character	as.c	-	LEAF	43	8.237548
as.list	as.l	-	-	0	0.000000
as.list.data.frame	a...	-	LEAF	22	4.214559
as.list.default	as..	-	LEAF	1	0.191571
as.name	as.n	-	LEAF	1	0.191571
coef	coef	-	LEAF	1	0.191571
deparse	dprs	-	LEAF	1	0.191571
doTryCatch	dTrC	-	-	0	0.000000
eval	eval	-	LEAF	1	0.191571
evalFunc	evlF	-	-	0	0.000000
file	file	-	LEAF	1	0.191571
FUN	FUN	-	LEAF	1	0.191571
lapply	lppl	-	LEAF	2	0.383142
lazyLoadDBfetch	lLDB	-	LEAF	2	0.383142
list	list	-	LEAF	5	0.957854
lm	lm	-	LEAF	42	8.045977
lm.fit	lm.f	-	LEAF	87	16.666667
match	mtch	-	LEAF	1	0.191571
mean	mean	-	-	0	0.000000
mean.default	mn.d	-	LEAF	2	0.383142
mode	mode	-	LEAF	2	0.383142
model.frame	mdl.f	-	-	0	0.000000
model.frame.default	mdl.f.	-	LEAF	12	2.298851
model.matrix	mdl.m	-	-	0	0.000000
model.matrix.default	mdl.m.	-	LEAF	51	9.770115
model.response	mdl.r	-	LEAF	13	2.490421
na.omit	n.mt	-	LEAF	20	3.831418
na.omit.data.frame	n...	-	LEAF	26	4.980843
names	nams	-	LEAF	2	0.383142
NCOL	NCOL	-	LEAF	1	0.191571
paste	past	-	-	0	0.000000
pmatch	pmtc	-	LEAF	2	0.383142
rep.int	rp.n	-	LEAF	7	1.340996

sapply	sppl	- LEAF	1	0.191571
simplify2array	smp2	- -	0	0.000000
structure	strc	- LEAF	32	6.130268
summary	smmr	- -	0	0.000000
summary.lm	smm.	- LEAF	40	7.662835
Sweave	Swev ROOT	-	0	0.000000
terms	trms	- -	0	0.000000
terms.formula	trm.	- LEAF	1	0.191571
try	try	- -	0	0.000000
tryCatch	tryC	- -	0	0.000000
tryCatchList	trCL	- -	0	0.000000
tryCatchOne	trCO	- -	0	0.000000
unique	uniq	- LEAF	3	0.574713
unlist	unls	- -	0	0.000000
vapply	vppl	- LEAF	3	0.574713
withVisible	wthV	- -	0	0.000000
	total.time	total.pct		
!	2	0.383142		
..getNamespace	1	0.191571		
.deparseOpts	4	0.766284		
.getXlevels	26	4.980843		
[99	18.965517		
[/data.frame	99	18.965517		
[[8	1.532567		
[[/data.frame	8	1.532567		
%in%	4	0.766284		
<Anonymous>	522	100.000000		
\$	1	0.191571		
anyDuplicated	23	4.406130		
anyDuplicated.default	22	4.214559		
as.character	43	8.237548		
as.list	23	4.406130		
as.list.data.frame	22	4.214559		
as.list.default	1	0.191571		
as.name	1	0.191571		
coef	1	0.191571		
deparse	2	0.383142		
doTryCatch	521	99.808429		
eval	521	99.808429		
evalFunc	521	99.808429		
file	1	0.191571		
FUN	7	1.340996		
lapply	30	5.747126		
lazyLoadDBfetch	3	0.574713		
list	5	0.957854		
lm	474	90.804598		
lm.fit	113	21.647510		
match	11	2.107280		
mean	2	0.383142		
mean.default	2	0.383142		
mode	2	0.383142		
model.frame	168	32.183908		
model.frame.default	168	32.183908		

model.matrix	69	13.218391
model.matrix.default	69	13.218391
model.response	56	10.727969
na.omit	134	25.670498
na.omit.data.frame	114	21.839080
names	2	0.383142
NCOL	1	0.191571
paste	1	0.191571
pmatch	2	0.383142
rep.int	7	1.340996
sapply	14	2.681992
simplify2array	4	0.766284
structure	33	6.321839
summary	520	99.616858
summary.lm	45	8.620690
Sweave	522	100.000000
terms	2	0.383142
terms.formula	1	0.191571
try	521	99.808429
tryCatch	521	99.808429
tryCatchList	521	99.808429
tryCatchOne	521	99.808429
unique	4	0.766284
unlist	1	0.191571
vapply	23	4.406130
withVisible	521	99.808429

```

Input
#str(profile_nodes_rle, max.level=2, vec.len=3, nchar.max=40, list.len=6)
strx(sumsprof01)

```

```

Output
'data.frame':      62 obs. of  7 variables:
 $ shortname : Factor w/ 62 levels "!", "..gN", ".dp0", ...: 1 2 3 4 5 6 7
   8 ...
 $ root      : Factor w/ 2 levels "-", "ROOT": 1 1 1 1 1 1 1 ...
 $ leaf      : Factor w/ 2 levels "-", "LEAF": 2 1 2 1 1 2 1 2 ...
 $ self.time : num 2 0 2 0 0 57 0 1 ...
 $ self.pct  : num 0.383 0 0.383 0 ...
 $ total.time: num 2 1 4 26 99 99 8 8 ...
 $ total.pct : num 0.383 0.192 0.766 4.981 ...

```

```

Input
#str(sumsprof01, max.level=2, vec.len=3,
#    nchar.max=40, list.len=6,
#    width=70, strict.width="wrap")

```

The classical approach hides the work that has been done. Actually it breaks down the data to record items. This figure is not reported anywhere. In our case, it can be reconstructed. The profile data have 8456 words in 524 lines.

In our approach, we break down the information. Two lines of control information are split off. We have 522 lines of profile with 50 unique stacks, referencing 62 nodes.

Instead of reducing it to a summary, we keep the full information. Information is always kept on its original level.

On the profiles level, we know the sample interval length, and the id of the stack recorded. On the stack level, for each stack we have a reference count, with the sample interval lengths used as weights. This reference count is added up for each node in the stack to give the node timings.


Cheap statistics are collected as they come by. For example, from the stacks table it is cheap to identify root and leaf nodes, and this mark is propagated to the nodes table.

3.1. Graph package.


Input

```
#12 6
oldpar <- par(mfrow=c(1,2))
library(graph)
plotviz(as(adjacency(sprof01),"graphNEL"),
        main="graph layout sprof01", sub=sprof01$info$id)
plotviz(as(adjacency(sprof02),"graphNEL"),
        main="graph layout sprof02", sub=sprof02$info$id)
par(oldpar)
```

graph layout sprof01



graph layout sprof02



R is function based, and control structures in general are implemented as functions. In a graph representation, they appear as nodes, concentrating and seeding to unrelated paths. We can detect these on the stack level and replace them by surrogates, introducing new nodes.

ToDo: fix null name


Input

```
sprof04 <- readRprof("RprofsRegressionExp103.out"); sprof04$info$id <- "sprof04"
#sprof04$nodes$name[1] <- sprof04$nodes$name[2]
#sprof04$nodes$name[1] <- "<noop>"??"
```

ToDo: cut top levels

Input


```
#24 12
library(graph)
a04<-adjacency(sprof04)
rnames <- rownames(a04)
rnames[1]<-"noop";rownames(a04) <- rnames; colnames(a04) <- rnames;
a04NEL <- as(a04,"graphNEL")
nodeRenderInfo(a04NEL) <- list(shape="ellipse")
nodeRenderInfo(a04NEL) <- list(cex=0.4, shape="ellipse")
a04NEL <- layoutGraph(a04NEL)
renderGraph(a04NEL)
```



Input

```
#24 12 --ellipses are lost?? needs to be recovered
plotviz(a04NEL, main="graph layout sprof04 plot")
```

graph layout sprof04 plot



4. STANDARD OUTPUT

For a reference, here are complete outputs of the standard functions.

`sprof <- sprof01` *Input*

4.1. **Print.** We omit the (lengthy) print output here and just give the commands as a reference.

`print_nodes(sprof)` *Input*

`print_stacks(sprof)` *Input*

`print_profiles(sprof)` *Input*

The `print()` method for `sprof` objects concatenates these three functions.

4.2. **Summary.**

`summary_nodes(sprof)` *Input*

`summary_stacks(sprof)` *Input*

`summary_profiles(sprof)` *Input*

The `summary()` method for `sprof` objects concatenates these three functions.

ToDo: Clarify: "print prints its argument and returns it invisibly (via `invisible(x)`)."

Return the argument, or some print representation?

ToDo: is there a `print=FALSE` variant to postpone printing to e.g. `xtable`?

4.3. Plot.

5. GRAPH

Graph layout is a theme of its own. Proposals are readily available, as are their implementation. For some of them, there are R interfaces or re-implementations in R. Their usefulness in our context has to be explored, and the answers will vary with personal preferences.

For some graph layout packages we illustrate an interface here and show a sample result. We use the original profile data here. This is a nasty graph with some R stack peculiarities. The corresponding results for the trimmed profile data are shown in the next section section 5.2 on page 66. This is a more realistic example of the kind of graphs you will have to work with.

ToDo: by graph package: preferred input format?

ToDo: use attributes. Edge with should be easy.

ToDo: include information from stack connectivity.

5.1. **Example: regression.** In this section, we use the recent version of our example, *sprof02* for demonstration. You can re-run it, using your *sprof* data by modifying this instruction:

Input

```
sprof <- sprof02
```

To interface *sprof* to a graph handling package, *adjacency()* can extract the adjacency matrix from the profile.

There are various packages for finding a graph layout, and the choice is open to your preferences. The R packages for most of these are just wrapper

Input

```
sprofadj <- adjacency(sprof)
```


This is a format any graph package can handle (maybe).

5.1.1. *graph* package.

```
library(graph)
sprofadjNEL <- as(sprofadj, "graphNEL")
```

```
plotviz(sprofadjNEL, main=paste("graph layout\n", sprof$info$id))
#detach("package:graph")
```

graph layout sprof02 updated

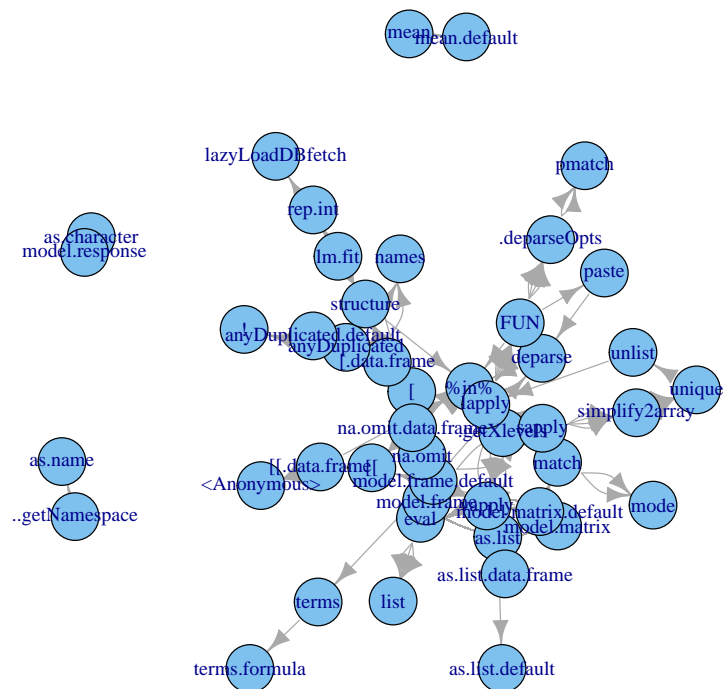


5.1.2. *igraph* package.

```
library(igraph)
sprof1 <- graph.adjacency(sprofadj)
```

```
plot(sprof1, main="sprof01: igraph layout")
#plot(sprof1, main="sprof01: igraph layout", cex.main=2)
#plotviz(sprof1, main="sprof01: igraph layout")

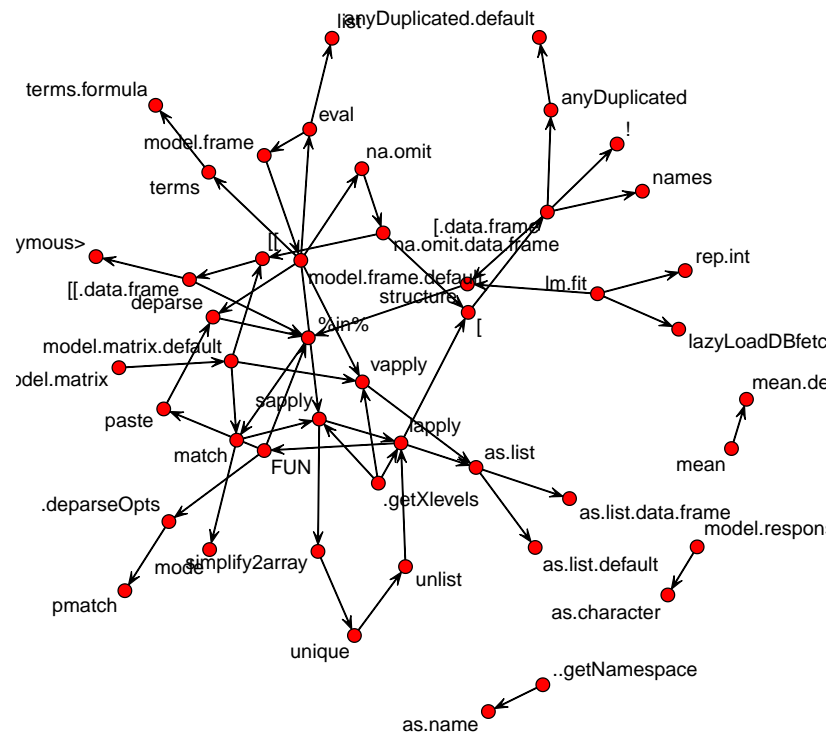
detach("package:igraph")
```

sprof01: igraph layout

5.1.3. *network package.*

```
library(network)
nwsprofadj <- as.network(sprofadj) # names is not imported
network.vertex.names(nwsprofadj) <- rownames(sprofadj) # not honoured by plot
plot(nwsprofadj, label=rownames(sprofadj), main="sprof01: network layout", cex.main=2)
```

sprof01: network layout




Experiments to include weight.

ToDo: maximum
edge.lwd?

Input

```
edge.lwd<-sprofadj
edge.lwd[edge.lwd>0]<- rank(edge.lwd[edge.lwd>0], ties.method="min")
#edge.lwd <- trunc(sprofadj/max(sprofadj)*10)+1
edge.lwd <- round(edge.lwd/max(edge.lwd)*12)
plot(nwsprofadj, label=rownames(sprofadj), main="sprof01: network layout", cex.main=2, edge.lwd=edge.lwd)
detach("package:network")
```

sprof01: network layout




5.1.4. *Rgraphviz* package.

```
library(Rgraphviz)
sprofadjRag <- agopen(sprofadjNEL, name="Rprof Example")
```


```
# 8 8                                     Input
plotviz(sprofadjRag, main="sprof01: Graphviz dot layout")
```

sprof01: Graphviz dot layout



#8


sprof01: Graphviz neato layout



#6 *Input*


```
plot(sprofadjRag,"twopi", main="sprof01: Graphviz twopi layout")
```

sprof01: Graphviz twopi layout




Input
`plot(sprofadjRag,"circo", main="sprof01: Graphviz circo layout")`

sprof04: Graphviz circo layout



Input
`plot(sprofadjRag,"fdp", main="sprof01: Graphviz fdp layout")`

sprof01: Graphviz fdp layout



5.2. Trimmed example: regression. In this section, we use the reduced version of our example, *sprof04* for demonstration. Except for the change of the data set, this is just a copy of the previous chapter, collecting the various layouts for easy reference.

Some experiments may have found their way to this chapter. They will be expelled.

You can re-run it, using your *sprof* data by modifying this instruction:

```
sprof <- sprof04
```

To interface *sprof* to a graph handling package, *until()* can extract the adjacency matrix from the profile.


```
sprofadj <- adjacency(sprof)  
adjname <- colnames(sprofadj)  
adjname[adjname==""] <- "<NULL>"  
colnames(sprofadj) <- adjname  
rownames(sprofadj) <- adjname
```

This is a format any graph package can handle (maybe).

5.2.1. *graph* package.


```
library(graph)
sprofadjNEL <- as(sprofadj, "graphNEL")
```

```
#8
plot(sprofadjNEL, main="sprof04: graph layout",)
#detach("package:graph")
```

sprof04: graph layout

```
#18
plot(sprofadjNEL, main="sprof04: graph layout",)
#detach("package:graph")
```


sprof04: graph layout



Input

```
#12
plot(sprofadjNEL, main="sprof04: graph layout")
#detach("package:graph")
```


sprof04: graph layout



5.2.2. *igraph* package.

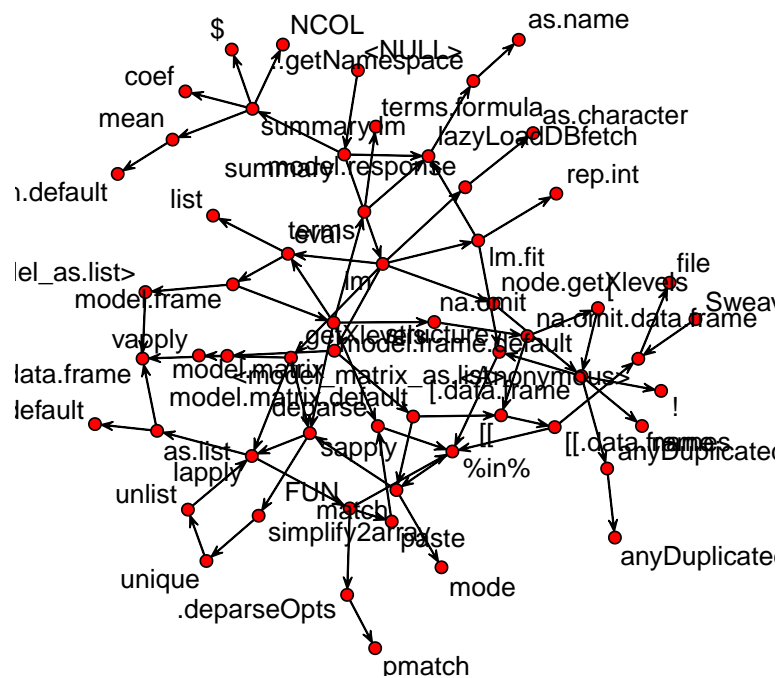
```
library(igraph)
sprof04 <- graph.adjacency(sprofadj)
```

```
#6
#plot(sprof04, main="sprof04: igraph layout", cex.main=2)
plot(sprof04, main="sprof04: igraph layout: trimmed data")
detach("package:igraph")
```

sprof04: igraph layout: trimmed data

5.2.3. *network package.*


```
#6
library(network)
nwsprofadj <- as.network(sprofadj) # names is not imported
network.vertex.names(nwsprofadj) <- rownames(sprofadj) # not honoured by plot
plot(nwsprofadj, label=rownames(sprofadj),
      main="sprof04: network layout: trimmed data", cex.main=2)
```

sprof04: network layout: trimmed dat


Input

```
#8
edge.lwd<-sprofadj
edge.lwd[edge.lwd>0]<- rank(edge.lwd[edge.lwd>0], ties.method="max")
#edge.lwd <- trunc(sprofadj/max(sprofadj)*10)+1
edge.lwd <- round(edge.lwd/max(edge.lwd)*12)
plot(nwsprofadj, label=rownames(sprofadj),
     main="sprof04: network layout: trimmed data",
     cex.main=2, edge.lwd=edge.lwd)
```

sprof04: network layout: trimmed data



#12


**sprof04: network kamadakawai layout:
trimmed data**

```

Input
plot(nwsprofadj, label=rownames(sprofadj),
     main="sprof04: network circle layout: \n trimmed data",
     mode="circle",
     cex.main=2, edge.lwd=edge.lwd)

```

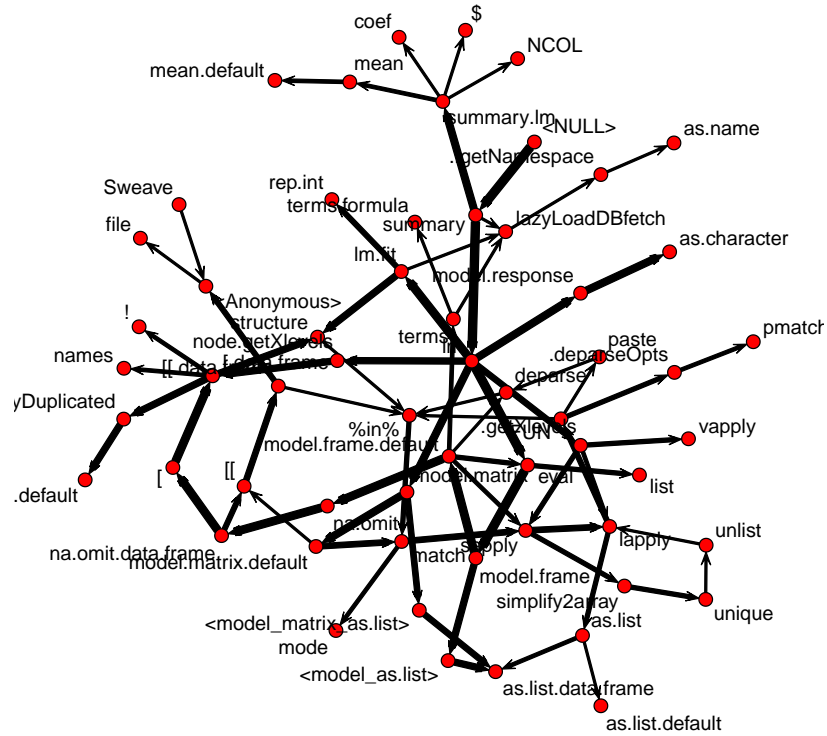
sprof04: network circle layout: trimmed data



```

Input
plot(nwsprofadj, label=rownames(sprofadj),
     main="sprof04: network fruchtermanreingold layout: \n trimmed data",
     mode="fruchtermanreingold",
     cex.main=2, edge.lwd=edge.lwd)
detach("package:network")

```


**sprof04: network fruchtermanreingold layout:
trimmed data**

5.2.4. *Rgraphviz* package.

```
library(Rgraphviz)
sprofadjRag <- agopen(sprofadjNEL, name="Rprof Example")
```

```
#6
plotviz(sprofadjRag, main="sprof04: Graphviz dot layout")
```


sprof04: Graphviz dot layout



#6 *Input*

```
plotviz(sprofadjRag,"neato", main="sprof04: Graphviz neato layout")
```

sprof04: Graphviz neato layout




Input

#6


```
plotviz(sprofadjRag,"twopi", main="sprof04: Graphviz twopi layout")
```

sprof04: Graphviz twopi layout




```
#6  
plotviz(sprofadjRag,"circo", main="sprof04: Graphviz circo layout")
```

sprof04: Graphviz circo layout



```
#6  
plotviz(sprofadjRag,"fdp", main="sprof04: Graphviz fdp layout")
```

sprof04: Graphviz fdp layout



INDEX

- Topic **adjacency**
 until, 66
- Topic **hplot**
 plot_nodes, 12
- Topic **manip**
 adjacency, 21
 list.as.matrix, 21
 profiles_matrix, 21
 stacks_matrix, 21
 stackstoadj, 21
- Topic **misc**
 apply, 5
 print, 55
 sampleRprof, 9
 summary, 55
 summaryRprof, 5
- Topic **util**
 adjacency, 56
- ToDo
- 0: add keep3 to keep header, some middle, tail, 3
 - 0: remove text vdots from string/name columns. Use empty string., 4
 - 1: Can we calibrate times to CPU rate? Introduce cpu clock cycle as a time base, 6
 - 1: Defaults by class, 16
 - 1: Move class attributes to package code, 14
 - 1: add a reference to colorbrewer, 16
 - 1: add class by keyword, 14
 - 1: add sampling.interval, sampling.time for backward compatibility, 9
 - 1: apply colour to selection?, 12
 - 1: classes need separate colour palette, 16
 - 1: colour by class – redo. Bundle colour index with colour?, 13
 - 1: improve colour: support colour in a structure, 12
 - 1: rearrange stacks? detect order?, 5
 - 1: remove text vdots from string/name columns, 12
 - 1: spread colour on displayed part, 12
 - 2: Implement. Currently best handled on source=text level, 35
 - 2: This section needs to be reworked, 30
 - 2: Warning: data structure still under discussion, 38
 - 2: add a purge function, 27
 - 2: add attributes to stacks, and discuss scope, 20
 - 2: add current level, 44
 - 2: add marginals and conditionals. Provide function node.summary., 40
 - 2: add names for node dimension, 39
 - 2: add time per call information: add marginals statistics run time by node, 44
 - 2: add trim by keyword, 30
 - 2: allow sorting, e.g. by marginals, 41
 - 2: check and stabilise colour linking, 21
 - 2: check and synchronise, 23
 - 2: colours. recolour. Propagate colour to graph., 25
 - 2: complete matrix conversion, 21
 - 2: cut next level, 34
 - 2: example, 24
 - 2: function addnode using “call by reference” to be added, 35
 - 2: generate a coplot representation, 44
 - 2: handle empty stacks and zero counts gracefully, 27
 - 2: implement, 36
 - 2: implement replacement on the stack level., 36
 - 2: keep as factor. This is a sparse cube with margins node, stack level, run length. Nodes are mostly concentrated on few levels., 38
 - 2: needs serious revision, 36
 - 2: re-think: sort stacks, 20
 - 2: replace by decent vector/array based implementation, 39
 - 2: replace sum by weighted sum, 41
 - 2: rescale to application scale, 41
 - 2: some mess in the code below – not working, 39
 - 2: sorting/arranging stacks, 20
 - 2: support na removal in rrle, 36
 - 2: table: node #runs min median run length max, 44
 - 2: trimexample, 30
 - 2: updateRprof needs careful checking. For now, we are including long listings here to provide the necessary information, 26
 - 2: use sprof02 or sprof03?, 36
 - 3: cut top levels, 53
 - 3: fix null name, 53
 - 3: use stack colours, 47
 - 4: Clarify:print prints its argument and returns it invisibly (via invisible(x)). Return the argument, or some print representation?, 55
 - 4: is there a print=FALSE variant to postpone printing to e.g. xtable?, 55
 - 5: by graph package: preferred input format?, 56
 - 5: include information from stack connectivity., 56
 - 5: maximum edge.lwd?, 60

5: use attributes. Edge with should be
easy., 56

adjacency, 21, 56

apply, 5

Index01, 80

list.as.matrix, 21

plot_nodes, 12

print, 55

profiles_matrix, 21

sampleRprof, 9

stacks_matrix, 21

stackstoadj, 21

summary, 55

summaryRprof, 5

until, 66

R session info:

- R version 3.0.1 (2013-05-16), x86_64-apple-darwin10.8.0
- Locale: en_GB.UTF-8/en_GB.UTF-8/en_GB.UTF-8/C/en_GB.UTF-8/en_GB.UTF-8
- Base packages: base, datasets, graphics, grDevices, grid, methods, stats, utils
- Other packages: graph 1.38.2, RColorBrewer 1.0-5, Rcpp 0.10.3, Rgraphviz 2.4.0, sna 2.3-1, sprof 0.0-5, wordcloud 2.4, xtable 1.7-1
- Loaded via a namespace (and not attached): BiocGenerics 0.6.0, igraph 0.6.5-2, network 1.7.2, parallel 3.0.1, slam 0.1-28, stats4 3.0.1, tools 3.0.1

L^AT_EX information: textwidth: 4.9823in linewidth:4.9823in

textheight: 8.0824in

Svn repository information:

```
$HeadURL: svnssh://gsawitzki@svn.r-forge.r-project.org/svnroot/sintro/pkg/sprof/vignettes/sprofiling.Rnw +
$Source: /u/math/j40/cvsroot/lectures/src/insider/profile/Rnw/profile.Rnw,v $
$Id: sprofiling.Rnw 187 2013-07-29 19:50:43Z gsawitzki $
$Revision: 187 $
$Date: 2013-07-29 21:50:43 0200(Mon, 29 Jul 2013)+
$name: $
$Author: gsawitzki $
```

GÜNTHER SAWITZKI
STATLAB HEIDELBERG
IM NEUENHEIMER FELD 294
D 69120 HEIDELBERG

E-mail address: gs@statlab.uni-heidelberg.de

URL: <http://sintro.r-forge.r-project.org/>